
NeuroMiner Manual

*PRONIA, Personalized Prognostic Tools for
Early Psychosis Management
Ludwig-Maximilians-Universität, München
Software written by Nikolaos Koutsouleris
Manual written by Dom Dwyer
Formatting, testing, and additional writing:
Carlos Cabral, Shalaila Haas, Anne Ruef, Adyasha Khuntia
Date: 6th September 2020
Version 1.4*

Contents

1	Introduction	1
2	Suggested Prerequisites	4
3	Getting Started	6
3.1	Compatibility	6
3.2	Downloading and Initialising NeuroMiner	6
3.3	Nomenclature	6
4	Main Interface	9
4.1	Input data into NM	13
4.2	Define parameter template	22
4.3	Multi-group analyses	57
4.4	Stacking	58
4.5	Initialize analyses	59
4.6	Preprocess features	62
4.7	Train supervised classifiers	63
4.8	Visualize classifiers	64
4.9	Display training results	66
4.10	Generate master files	74
4.11	Out of Sample Cross-Validation	74
4.12	Load NeuroMiner structure	78
4.13	Save NeuroMiner structure	78

4.14 Clear NeuroMiner structure	78
4.15 Change Working Directory	78
4.16 Open Manual	78
4.17 Utilities	78
5 Example	80
6 Known Issues	88

1 Introduction

Machine learning techniques are poised to become clinically useful methods that may be used for diagnosis, prognosis, and treatment decisions. Despite this, they are currently underutilised in medical studies and, even more in psychiatric research because most current tools require strong programming and computational engineering skills (e.g., scikit-learn, caret, Weka, nilearn). While there are some great tools that do not require programming experience (e.g., PRoNTo), these are often focused on making predictions from specific data domains such as neuroimaging data. This highlights a pressing need for user-friendly machine learning software that makes advanced methods available to clinical researchers from different fields aiming at collaboratively developing diagnostic, predictive, and prognostic tools for precision medicine approaches.

NeuroMiner has been continuously developed by Nikolaos Koutsouleris since 2009 to provide clinical researchers with cutting-edge machine learning methods for the analysis of heterogeneous data domains, such as clinical and neuropsychological read-outs, structural and functional neuroimaging data, and genetic information. The program can be considered as an interface to a large variety of unsupervised, and supervised pattern recognition algorithms that have been developed in the machine learning field over the last decades. Furthermore, the application implements itself different strategies for preprocessing, filtering and fusing heterogeneous data, training ensembles of predictors and visualizing and testing the significance of the computed predictive patterns. The current release candidate of NeuroMiner has been tested in the Section of Neurodiagnostic Applications on a variety of datasets from healthy controls and patients with psychiatric disorders and was designed specifically to create robust models with a high probability of generalization to new datasets. For reference, we include here a list of papers (see below and publications listed in [PubMed](#)), which were all based on previous versions of the program.

More specifically, using a light-weight and interactive text-based menu system, NeuroMiner allows the user to:

- a) load their data easily (e.g., using spreadsheets, NifTi images, or SPM structures);
- b) build a variety of cross-validation frameworks for classification and regression problems that have become a gold standard in the field (e.g., repeated nested cross-validation, leave-site-out cross-validation);
- c) apply a range of preprocessing strategies (e.g., scaling, filtering, many forms of dimensionality reduction, etc.);
- d) choose and combine cutting-edge supervised algorithms (e.g., support vector machine, elastic net, random forest, etc.);
- e) apply feature selection procedures (e.g., wrappers), data fusion techniques, and stacked generalization;
- f) apply learned models to new data (external validation).

To assist in selecting and analysing data, the user can visualise the data during input, monitor accuracy during learning, and understand the results of complex analyses using multiple display options. These allow the user to accurately report the data and also to understand the underlying machine learning analyses. Furthermore, the ability to apply the learned models to completely new data is important because it is quickly becoming a standard requirement of all machine learning studies. Combined, NeuroMiner gives the user the opportunity to design, implement, understand, and report machine learning analyses.

DISCLAIMER

Please note that NeuroMiner is supplied as is and no formal maintenance is provided or implied. In no event shall the author of the software (heretofore known as the Author) be liable to any party for direct, indirect, special, incidental, or consequential damages, including lost profits, arising out of the use of this software and its documentation, even if the Author has been advised of the possibility of such damage. The Author specifically disclaims any warranties, including, but not limited to, the implied warranties of merchantability and fitness for a particular purpose. The software and accompanying documentation provided hereunder is provided 'as is'. The Author has no obligation to provide maintenance, support, updates, enhancements, or modifications (but we plan to).

This is the beta release version of the software and the software is undergoing regular updates. Please send any comments, questions, or bug reports to email.neurominer@gmail.com.

Papers that have used used NM:

1. Koutsouleris N, Meisenzahl EM, Davatzikos C, Bottlender R, Frodl T, Scheuerecker J, Schmitt G, Zetsche T, Decker P, Reiser M, Moller HJ, Gaser C. Use of neuroanatomical pattern classification to identify subjects in at-risk mental states of psychosis and predict disease transition. *Archives of General Psychiatry*. 2009; 66(7):700-12.
2. Koutsouleris N, Meisenzahl EM, Borgwardt S, Riecher-Rossler A, Frodl T, Kambeitz J, Kohler Y, Falkai P, Moller H.-J., Reiser M, Davatzikos C. Individualized differential diagnosis of schizophrenia and mood disorders using neuroanatomical biomarkers. *Brain*. 2015 Jul;138(Pt 7):2059-73.
3. Koutsouleris N, Kahn RS, Chekroud AM, Leucht S, Falkai P, Wobrock T, Derkx EM, Fleischhacker WW, Hasan A. Multisite prediction of 4-week and 52-week treatment outcomes in patients with first-episode psychosis: a machine learning approach. *Lancet Psychiatry*. 2016 Oct;3(10):935-946. doi: 10.1016/S2215-0366(16)30171-7.
4. Cabral C, Kambeitz-Illankovic L, Kambeitz J, Calhoun VD, Dwyer DB, von Saldern S, Urquijo MF, Falkai P, Koutsouleris N. Classifying Schizophrenia Using Multimodal Multivariate Pattern Recognition Analysis: Evaluating the Impact of Individual Clinical Profiles on the Neurodiagnostic Performance. *Schizophrenia Bulletin*. 2016 Jul;42 Suppl 1:S110-7. doi: 10.1093/schbul/sbw053.

5. Koutsouleris N, Borgwardt S, Meisenzahl EM, Bottlender R, Moller HJ, Riecher-Rossler A. Disease prediction in the at-risk mental state for psychosis using neuroanatomical biomarkers: results from the FePsy-study. *Schizophrenia Bulletin*. 2012; 38(6):1234-46
6. Borgwardt SJ, Koutsouleris N, Aston J, Studerus E, Smieskova R, Riecher-Rossler A, Meisenzahl EM. Distinguishing prodromal from first-episode psychosis using neuroanatomical pattern recognition: Evidence from single-subject structural MRI. *Schizophrenia Bulletin*. 2013; 39(5):1105-14. doi: 10.1093/schbul/sbs095
7. Koutsouleris N, Davatzikos C, Bottlender R, Patschurek-Kliche K, Scheuerecker J, Decker P, Gaser C, Moller HJ; Meisenzahl E. Early recognition and disease prediction in the at-risk mental states for psychosis using neurocognitive pattern classification. *Schizophrenia Bulletin*. 2012; 38(6):1200-15
8. Koutsouleris N, Riecher-Rossler A, Meisenzahl E, Smieskova R, Studerus E, Kambeitz-Ilankovic L, von Saldern S, Cabral C, Reiser M, Falkai P, Borgwardt S. Detecting the psychosis prodrome across high-risk populations using neuroanatomical biomarkers. *Schizophrenia Bulletin*. 2014, 41(2):471-82.
9. Koutsouleris N, Davatzikos C, Borgwardt S, Gaser C, Bottlender R, Frodl T, Falkai P, Riecher-Rossler A, Moller HJ, Reiser M, Pantelis C, Meisenzahl E. Accelerated Brain Aging in Schizophrenia and Beyond: A Neuroanatomical Marker of Psychiatric Disorders. *Schizophrenia Bulletin*. 2014 Sep;40(5):1140-53
10. Koutsouleris N, Gaser C, Bottlender R, Davatzikos C, Decker P, Jager M, Schmitt G, Reiser M, Moller HJ, Meisenzahl EM, Use of Neuroanatomical Pattern Regression to Predict the Structural Brain Dynamics of Vulnerability and Transition to Psychosis. *Schizophrenia Research*. 2010;123(2-3):175-187
11. Kambeitz-Ilankovic L, Meisenzahl EM, Cabral C, von Saldern S, Kambeitz J, Falkai P, Moller HJ, Reiser M, Koutsouleris N. Prediction of outcome in the psychosis prodrome using neuroanatomical pattern classification. *Schizophrenia Research*. 2015;173(3):159-65.
12. Koutsouleris et al., *JAMA Psychiatry*. Prediction models of functional outcomes for individuals in the clinical high-risk state for psychosis or with recent-onset depression: a multimodal, multisite machine learning analysis. 2018.

2 Suggested Prerequisites

NeuroMiner works exclusively from the MATLAB command line and at some points the variables need to be in specific formats for NeuroMiner to recognise them, as described later in this manual. You may also want to investigate the NeuroMiner outputs stored in output variables. For these reasons, we recommend a minimal level of knowledge about command-line operation of MATLAB (see below).

MATLAB:

- Differences between vectors/matrices/cell-arrays/structures
- How to address certain cells/rows/columns of them
- Deleting /creating/addressing items in a matlab structure
- Differences between numerical/string/binary variables
- How to add stuff to the MATLAB path and why this is needed
- How to import files into MATLAB

→ For basic tutorial

<http://antoniahhamilton.com/matlab.html>

NeuroMiner is also software that uses a number of advanced machine learning tools. We have provided some description of the most important concepts and techniques, but the manual often assumes machine learning knowledge because it was not possible to create a full textbook. For this reason, we recommend a minimal knowledge of machine learning methods (see below).

Machine learning:

- Difference between training and test samples?
- What is crossvalidation/nested crossvalidation?
- What is preprocessing (for machine learning, not just MRI data)?
- What are filters and wrappers?
- What are features and classification targets?
- What is Principal Component Analysis (PCA)?
- What is a Support Vector Machine (SVM)?

→ For basic tutorial

<http://www.autonlab.org/tutorials/>

→ Suggested literature

The following literature is recommended for a basis of knowledge prior to use of the program:

- <http://www.ncbi.nlm.nih.gov/pubmed/?term=Machine+learning+classifiers+and+fMRI%3A+a+tutorial+overview>
- <http://www.sciencedirect.com/science/article/pii/S2213158213001204>
- <http://journal.frontiersin.org/article/10.3389/neuro.09.032.2009/full>

- <http://journal.frontiersin.org/Journal/10.3389/neuro.09.032.2009/full>
- <http://www.ncbi.nlm.nih.gov/pubmed/?term=Why+voxel-based+morphometric+analysis+should+be+used+with+great>

3 Getting Started

3.1 Compatibility

NeuroMiner has been tested using Windows, Linux (Cent OS), and Mac OS operating systems and works with MATLAB2018b and above. We recommend the "Statistics and Machine Learning" and the "Optimization" toolboxes in order to use advanced training options in the matLearn toolbox and for imputation functions.

For non-supported operating systems, NeuroMiner has a number of programs installed that may require independent compilation using a C compiler (e.g., gcc for linux or Xcode for mac).

3.2 Downloading and Initialising NeuroMiner

NeuroMiner works in two main modes: matrix and neuroimaging (i.e., NIFTI or FreeSurfer files). For neuroimaging using NIFTI data format, it is necessary to have a working copy of the Statistical Parametric Mapping (SPM) toolbox installed ([download](#)). We also recommend downloading the WFU Pickatlas toolbox ([link](#)) to take advantage of this selection tool when inputting data. For surface-based analyses, it is necessary to have a FreeSurfer distribution installed ([download](#)).

Once the program is downloaded, it needs to be added to the MATLAB path by using the **addpath** command from the command line or by using the "Set Path" button on the toolbar (e.g., "addpath /home/NM/NeuroMiner"). NeuroMiner can then be launched by typing **nm** on the command line. To load the program faster, you can type **nm nosplash** and to enter expert mode type **nm nosplash expert**. The first time that NeuroMiner is run, a file selector box will appear first asking to define the SPM directory and then the FreeSurfer directory. If these directories are not available, then simply press cancel and NeuroMiner will be launched in non-imaging mode—i.e., a mode that restricts options only to matrix data (see Fig. 1; upper). Otherwise, NeuroMiner will be launched in matrix and neuroimaging mode (Fig. 1; lower).

User note: once NeuroMiner is initialised and added to the path, there can be incompatibilities between some MATLAB functions on some systems. For example, there can be a incompatibility with the function **ttest2**. In these cases, NeuroMiner needs to be removed from the path in order to use the original function.

3.3 Nomenclature

The field of machine learning uses different nomenclature to classical statistics. The following terms and their definitions are used throughout the manual (Rob Tibshirani, Statistics 315a, Stanford):

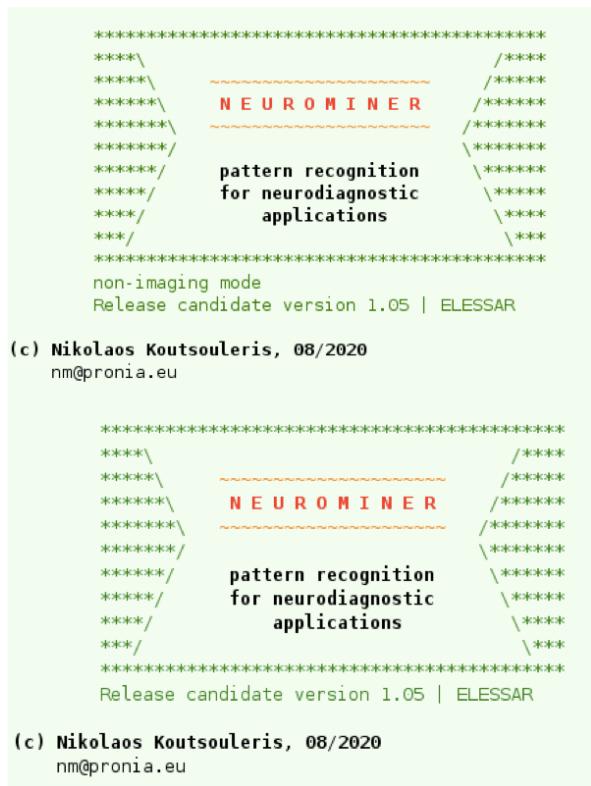


Figure 1: Display screen of NeuroMiner in different modes. NeuroMiner will be loaded in non-imaging mode (upper) when the SPM or FreeSurfer directories are not provided during start-up (upper). If the directories have been provided it will be started in standard imaging and matrix data mode (lower).

Machine Learning	Statistics
Features	Predictor variables (regression)
Target variable/labels	Outcome Variables
Weights	Parameters
Hyperparameters	Parameters
Learning	Fitting
Generalization	Test set performance
Supervised learning	Regression/Classification
Unsupervised learning	Clustering,density estimation
Large grant = \$1,000,000	large grant = \$50,000

```

*****
***\   NEUROMINER  *****
***\-----/*****
***\-----/*****
*****| pattern recognition  \*****
*****| for neurodiagnostic \*****
*****| applications \*****
***| \*****
***| \*****
*****| Release candidate version 1.05 | ELESSAR
(c) Nikolaos Koutsouleris, 08/2020
nm@pronia.eu

Current working directory: /home/dominic/NeuroMiner_Elessar_1.05
=====
MAIN INTERFACE [ DATA INPUT MODE ]
=====
1 | Load data for model discovery and cross-validation
2 | Load NeuroMiner structure
3 | Change working directory
4 | Utilities
<==| Back/Quit [Q]

Menu choice (1-4/Q) (Default: Q) ?

```

Figure 2: The main starting display of NeuroMiner

4 Main Interface

NeuroMiner is software that has a text-based interface and works from the MATLAB command line. This implies that all analyses are implemented by selecting options from text menus and entering parameters when asked to do so. Each of the menus share a similar format as described in the following example.

After the neuroimaging paths have been established, the first time that NeuroMiner is run the user will see the **MAIN INTERFACE** (see Fig.2) that consists of 4 different options that are numbered 1-4. Below these numerical options there is also an option to go **Back/Quit (Q)**. Below the dividing line (i.e., =====), the user can enter their choice of option in the space after **Menu choice (1-4/Q) (Default: Q)?**.

In this example, the user can do one of three things:

- 1) to select one of the 4 numerals corresponding to each option and press the enter key on the keyboard;
- 2) to enter "Q" to go back to the previous menu, or if they are in the **MAIN INTERFACE** to quit the program;
- 3) to simply press the enter key and the program will activate the default option—in this case, to quit the program.

For example, a user could press "2" followed by the enter key to "Load NeuroMiner structure". The NeuroMiner MATLAB structure ("NM") is the primary variable that NeuroMiner uses to store data, settings, analyses, and results. For example, when a user inputs data into NeuroMiner it is stored in the NM

structure as a data matrix. Then when the user chooses settings or performs an analysis, all of these inputs and outputs are stored in the structure as well. This means that the NM structure forms the core of any analysis within NeuroMiner.

→ **Critical Point**

It is important to note that the NeuroMiner structure is not automatically saved during analyses in order to give the user more control. This means that the user must save the structure periodically during the analysis using the menu item **10: Save NeuroMiner structure**.

NeuroMiner will **automatically change** the items in all menus based on two main criteria:

- 1) Whether a NeuroMiner MATLAB structure storing all the information from a previous analysis has been loaded to the MATLAB workspace;
- 2) The individual options that are selected during a NeuroMiner session.

For example, if a NM structure is loaded prior to the start of a NeuroMiner session then the program will automatically detect this and the **MAIN INTERFACE** will specifically relate to the loaded analysis. Similarly, if specific options are chosen during the establishment of the analysis then the menus will adapt to the chosen settings and options that do not apply to the analysis will be removed.

The Main Interface After Data Has Been Loaded and Processed

Once the user has entered and processed data (see section [4.1](#)), the **MAIN INTERFACE** will change to the following:

- 1: Inspect data used for model discovery and cross-validation
- 2: Set up NM parameter workspace
- 3: Initialize & manage analyses
- 4: Preprocess features
- 5: Train supervised classifiers
- 6: Visualise classifiers
- 7: Display training results
- 8: Generate master files
- 9: Load NeuroMiner structure
- 10: Save NeuroMiner structure
- 11: Clear NeuroMiner structure
- 12: Change working directory
- 13: Open Manual
- 14: Utilities
- Back/Quit[Q]

These options are major functional domains related to the methods and stages of machine learning analysis using NeuroMiner. The contents of each of these options are described throughout the manual, but it is important to understand the overall analytic workflow that the program uses by briefly describing each

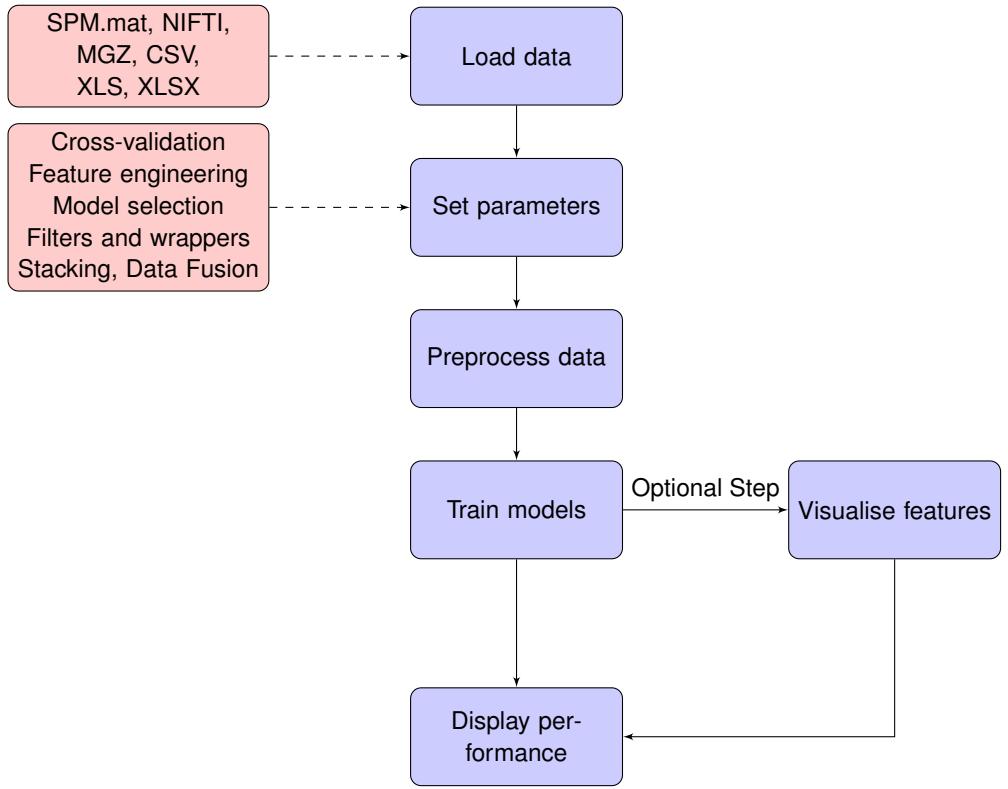


Figure 3: A flow chart for NeuroMiner. The main processing steps involve entering data, then defining a parameter template with the settings that are required for an analysis. It is then recommended that the data are preprocessed (i.e., that feature engineering takes place) and then the models can be trained. The user then has the option to 'visualize' the classifiers, where the model and other statistics on each individual feature are calculated. The user then can display the performance of the models, which will lead to a results screen.

of the options above.

1: Input data into NM The entry of data is the most important aspect of any analysis. When data is entered into NeuroMiner, it is stored within a MATLAB structure called "NM" (i.e., "NeuroMiner") in the workspace. All other analyses and results are then added to this NM structure as the user progresses. The contents of the structure can also be viewed from the MATLAB command line.

2: Define parameter template Once the data have been entered, the first step of any analysis is to define machine learning options. Within this menu, the user can establish the cross-validation framework (e.g., nested cross-validation), perform essential preprocessing operations on the data (e.g., scaling), choose the machine learning algorithm (e.g., SVM), define the machine learning algorithm settings (e.g., C values), and perform other operations on the data.

3: Initialize/Delete analyses After the entry of all the settings for an analysis, the user can then 'initialize' these settings in preparation for processing. This option saves the parameters in an analysis structure, which the user can then

select when they go to process the data. This is a great feature because it means that the user can define multiple different analyses for the same dataset by changing parameters in the 'parameter template' and then initializing separate analyses.

4: Preprocess Features In NeuroMiner, the analysis can be split into two parts: preprocessing features and training supervised classifiers. Preprocessing in NeuroMiner can also be called 'feature preparation and selection', and involves preparing the data for machine learning analysis (e.g., a support vector machine). For example, the researcher may want to covary for age and sex, perform a Principal Components Analysis, and scale neuroimaging data prior to analysis. Neurominer performs all preprocessing steps within the cross-validation structure, which is recognised as the most valid approach. Allowing the user to preprocess their data first within the cross-validation structure can save a lot of computational resources and time because it means that the researcher can run multiple analyses from the same preprocessed data.

5: Train supervised classifiers This option allows the user to select and run the machine learning analysis. For example, the user can run a support vector machine across a range of C parameters. The results of this analysis are then stored in the NM structure.

6: Visualise classifiers This step is optional. In a standard machine learning analysis, it is often unclear how each of the features (e.g., imaging voxels or variables) contribute to the prediction accuracy and they are often not in the same space as the original images. For this reason, the user has the option within NeuroMiner calculate statistics for each feature with the 'visualise classifiers' option. When using neuroimaging data, this option also projects (i.e., transforms) the data into the original space so that it can be viewed in a similar manner as other neuroimaging results in nii files.

7: Display training results Based on the above essential steps, and the optional step to visualise the classifiers, the results of the analysis will be ready to be viewed with this step. For example, the user can view the accuracy, the frequency of model selection, the classifier weights that are applied to each variable, and any additional visualisation routines that have been run.

8: Generate master files As mentioned above, NeuroMiner stores settings and analysis results in the NM structure within MATLAB, which is visible in the workspace. In addition to this, NeuroMiner can also store individual files related to steps **4: Preprocess Features** and **5: Train supervised classifiers** above. When NeuroMiner does this, it stores a single file for each permutation within the cross-validation framework—as described later. These individual files can then be saved as "master files" using this option.

9 - 11: Load, save, and clear NeuroMiner structure These options allow the user to load the NM structure (where all the settings and analyses are stored) to the MATLAB workspace, to save the structure as a file, and to clear the

structure from the workspace.

12: Change working directory In similarity to most neuroimaging software, NeuroMiner operates from a "working directory" where preprocessing or analysis files are saved during any analysis. The working directory is automatically the directory where NeuroMiner is started, but can be changed using this item.

13: Open manual This opens the manual that you are reading right now.

14: Utilities Contains functions related to FreeSurfer and ability to add SPM/FreeSurfer paths to the NeuroMiner environment.

4.1 Input data into NM

NeuroMiner works with 3D neuroimaging (i.e., NIFTI, SPM, or FreeSurfer) and/or numerical data stored in matrices (e.g., Excel, csv, or MATLAB formats). For example, these are some of the data you can enter:

- structural neuroimaging data (e.g., NIFTI files from an SPM study)
- functional neuroimaging data (e.g., NIFTI files from 2nd level contrasts)
- FreeSurfer data (e.g., thickness or gyration)
- clinical or cognitive data (e.g., in Excel)
- genetics or other biological data (e.g., in csv)
- functional connectivity matrices (e.g., in the MATLAB workspace)
- DTI streamline matrices (e.g., in the MATLAB workspace)

In NeuroMiner, each dataset is called a 'modality' and multiple modalities can be entered at the same time to allow data fusion (discussed below and in section 4.2.1). For example, you can choose to perform machine learning on structural neuroimaging data alone or in combination with cognitive data.

If data has not been entered and NeuroMiner does not detect a "NM" structure in the MATLAB workspace, then the user will be given the option to "1: load data for model discovery and cross-validation". Selecting this option will transfer the user to the "NM data workspace manager", which is an adaptive menu that changes based on selections, such as the machine learning framework or the data input format.

If no data has been selected the NeuroMiner menu will display a default menu that is configured for classification using NIFTI (neuroimaging) :

- 1 : Define machine learning framework [classification]
- 2 : Select data input format [nifti]
- 3 : Define no. of samples and sample identifiers [???]
- 4 : Involve unlabeled cases in the analyses [no]
- 5 : Select space-defining image [?]
- 6 : Read in additional image(s) for weighting [no]
- 6 : Describe data [?]

The machine learning framework can be changed to regression, which will modify some of the input entries that should be clear if the classification frame-

work is understood (e.g., changing the group labels to a continuous measure).

If you are in neuroimaging mode (i.e., if you have established paths for SPM and/or FreeSurfer), selecting the second option will change the menu to read:

- 1 : SPM (basic models)
- 2 : ANALYZE/NIFTI images (3D)
- 3 : Cortical surface files (3D)
- 4 : Matrix data (2D)

These options will now be outlined in order.

Box: Notification, Warning, & Error Messages. When data is entered, NeuroMiner checks whether it has been entered successfully and performs a number of pre-screening quality control measures. Messages related to successful data entry (**blue font**), unsuccessful data entry (**red font**), or warnings (**cyan font**) are issued at the top of the screen, such as:

"**Error in RetrievelImageInfo: Image /your/directory/here/image.img not found!**"
"**Warning in check matrix: I found only one column in the MATLAB workspace**"
"**Label tests passed**" or "**Predictor matrix tests passed**"

Please take note of these warnings and errors.

4.1.1 SPM

NeuroMiner gives the option of simply entering in an SPM.mat structure from a previous univariate analysis (e.g., t-tests, ANOVA, or regression). Each of the options in the following menu need to be selected:

- Define machine learning framework [regression]
- Select data input format [spm]
- Select path to SPM.mat [?]
- Review SPM design
- Specify which columns to extract as labels from SPM design [?]
- Select space-defining image [?]
- Read in additional image(s) for weighting
- Describe data [?]

Select path to SPM.mat A pop-up dialogue box will appear where you need to select the SPM.mat structure.

Review SPM design This will display the design using SPM tools.

Specify which columns to extract as labels from SPM design Labels represent what you want to predict; e.g., groups or regression targets. These need to be included in the SPM.mat structure and the column corresponding to the label needs to be selected here; e.g., "1" for the first column, or "1:2" for the first and second columns (e.g., if they are dummy-coded).

Select space-defining image Please see the description in the box "Space Defining Image" regarding selection of brain regions using the WFU PickAtlas.

Box: Space Defining Image

In NeuroMiner, a "space-defining image" refers to an image that will define two things: 1) the voxel dimensions; and 2) the mask used to select the data. This means that **all images will be resliced to the mask dimensions**.

NeuroMiner has two options for the mask: NeuroMiner has two options for the mask: 1) Select a mask using WFU Pickatlas if it is installed in your SPM path (see website); 2) use your own spacedefining image (e.g., the ICBM template). NeuroMiner will also give the option to choose WFU PickAtlas if the toolbox is installed following WFU PickAtlas directions (see [pdf](#)). Within WFU PickAtlas, the default option is to write out a combined mask of brain regions selected. However, if the option to "write independent regions" is selected (checkbox, bottom-middle) then the regions will be written as separate data modalities into NeuroMiner. This means that different algorithms could be used on the different regions (e.g., one linear and one non-linear).

Following selection of the mask, the user will be given the option to include voxel above a mask threshold (e.g., in the case of a binary mask, more than 0). This will be reflected within a option called 'Do you want to threshold the space-defining image?'. Setting the default options will allow the user to select for 'Define threshold(s)', 'Define threshold mode' and 'Define threshold operator' which will set a threshold for the space-defining image. If a mask has been created using common default settings within SPM/FSL, then it will not be a completely binary mask and there will be a gradient between 0-1 at the boundary of the mask. This means that a common selection for thresholding is more than 0.5. When using WFU PickAtlas and writing independent regions, this option should not be selected.

The brainmask will be stored so that if further external validation analyses are conducted then the following users do not need to have the original brainmask.

Read in additional image(s) for weighting This option gives you the opportunity to enter an image that will be used during preprocessing to weight the brain image features. It is specifically designed for when parcellated brain regions have been entered or when the WFU PickAtlas has been used to select brain parcels. In these cases, the image will consist of independent brain regions (e.g., the hippocampus, frontal lobe, cingulate cortex etc.). You can therefore enter a map of the same brain regions that have weights associated with each parcel based on the degree of importance of each region; e.g., 1 for the hippocampus, 2 for the frontal lobe, 3 for the cingulate. This weight map can then be used during preprocessing to up- or down-weight the features in the analysis (see section [4.2.3.12](#)).

Describe data Enter a short description of this data for later reference if there is more than one modality; e.g., "cobre structural imaging data".

IMPORT Once the template is completed, you'll be given the option to import the data into the NeuroMiner structure (i.e., the "NM" structure) in the MATLAB workspace. The subject identifiers will be the unique fields in the paths to the

files (e.g., the subject directory or filename).

4.1.2 ANALYZE/NIFTI images

Analyze or nifti images can be selected here. Once this option is selected, the menu will change to the following:

- 1 : Define machine learning framework [classification]
- 2 : Select data input format [nifti]
- 3 : Define no. of samples and sample identifiers [???]
- 4 : Involve unlabeled cases in the analyses [no]
- 5 : Select space-defining image [?]
- 6 : Read in additional image(s) for weighting
- 7 : Describe data [?]

Define no. of samples and sample identifiers Enter the number of groups (e.g., "2") and then enter descriptors for these groups (e.g., "patients" or "controls"). It's important to note that sensitivity will be defined as the correct classification of individuals in the first group entered, so this is usually patients in clinical studies.

Involve unlabelled cases in the analysis This gives the option to use cases that do not have a label for semi-supervised learning.

Select space-defining image Please see the "Space Defining Image Box".

Read in additional image(s) for weighting An image can be entered here and it can be used in later processing steps to weight the entered feature images. This functionality was primarily introduced in order to enter

Map image files to samples Use the SPM dialogue box pop-up to select the images for the samples together or for the groups separately. We recommend to print the absolute paths to the images in a text file and then use the "ed" feature of the SPM dialogue box (bottom right) to enter the paths. This can be one for the full group, where you will be prompted to enter the groups one after the other, or for each group separately.

Define global multiplier For each subject, multiply the voxel values by a numeric value.

Adjust data for globals This option is primarily for structural neuroimaging data where it is necessary to adjust the images for the global brain volume (e.g., whole-brain volume or intra-cranial volume). Either a vector is required in the MATLAB workspace or a text file is required. It is **critical** to note that the values must be in the same order as the images for your analysis—i.e., the participant order **must** match because there is no subject matching procedure.

Inspect image information and check registration This option will allow the user to check the registration parameters to the template image and then dis-

play the images using the SPM check-reg function. This is useful to check whether the images are going to be imported properly.

Read in additional image(s) for weighting This option gives you the opportunity to enter an image that will be used during preprocessing to weight the brain image features. For example, when parcellated brain regions have been entered or when the WFU PickAtlas has been used to select brain parcels the image will consist of independent brain regions (e.g., the hippocampus, frontal lobe, cingulate cortex etc.). You can therefore enter a map of the same brain regions that have weights associated with each parcel based on the degree of importance of each region; e.g., 1 for the hippocampus, 2 for the frontal lobe, 3 for the cingulate. This weight map can then be used during preprocessing to up- or down-weight the features in the analysis (see section [4.2.3.12](#) and 'Rank / Weight Features'). Alternatively, reliability coefficient images that are calculated during multi-site studies (e.g., with travelling subjects) can be entered here.

Describe data Enter a short description of the data (e.g., "cobre mwrp1 images"). This is important if there are more than one modality.

IMPORT nifti After successful completion of settings for Steps 1-7, an option for 'IMPORT nifti' will appear. Selecting this option will import the nifti images into the NM structure based on the space-defining image that was selected and based on any correction for global brain volume. The subject identifiers will be the unique fields in the paths to the files (e.g., the subject directory or filename).

4.1.3 Cortical surface files

This option is for FreeSurfer .mgz files. The menu items will change to be the same as the above for NIFTI files.

Please note that once NeuroMiner has finished processing, there is no utility currently to visualize the classifier results in the mgz format.

4.1.4 Matrix data: MATLAB workspace

Matrix data is defined as any numeric data that is stored in a matrix file, such as an Excel file, CSV file, .mat file, or in the MATLAB workspace. The default menu is to enter variables from the MATLAB workspace as outlined in this section, but other choices can be made by selecting "Define data provenance". The following section will outline the options specific to entry of data from the MATLAB workspace.

```
Define machine learning framework [ classification ]
Select data input format [ matrix ]
Define data provenance [ MATLAB workspace ]
Enter name of matrix variable containing the predictor data [ ? ]
Enter name of label variable in MATLAB workspace [ ? ]
```

Enter name of case ID variable in MATLAB workspace [?]

Enter name of feature descriptor variable in MATLAB workspace [?]

Describe data [?]

Enter name of matrix variable (cell array of string) in MATLAB workspace

Enter the matrix containing the features (i.e., variables or predictors) that you want to use to predict the labels (i.e., the target or outcome variable). This variable must be a standard [n x m] MATLAB double (subjects = n; features = m). That is, the subjects must be entered in the rows and the features are the columns.

Enter name of label variable (cell array of strings) in MATLAB workspace

Labels are the targets of the prediction. For classification, these are group names or values (e.g., "control" and "patient") and **must** be entered in a cell array of strings (see Box: Data Types in NeuroMiner). We strongly discourage the use of numbers to represent groups (e.g., "1" and "2") due to display and visualisation later in processing.

Box: Data types in NeuroMiner

In NeuroMiner there are two types of MATLAB that are used: standard double matrix and a cell array of strings. A standard double is what is created by default in MATLAB. A cell array of strings is when you have independent strings (e.g., `char('CTRL')`) stored in a cell string array (e.g., `cellstr(['ctrl';'pat'])`).

A common situation is when you have a vector containing numbers representing groups of interest (e.g., `groups = [1; 1; 2; 1; 2]`). We strongly recommend changing this to have string variables representing the groups (i.e., "CTRL" and "PAT").

Enter the name of case ID variable (cell array of strings) in MATLAB workspace

Enter a vector with the case IDs of each subject. This must be a cell array of strings (see Box: Data Types in NeuroMiner).

Enter name of feature descriptor variable (cell array of strings) in MATLAB workspace

Enter a vector with the names of the features. This must be a cell array of strings.

Inspect matrix data and select features for import Once the previous settings are entered, this option appears which will open a results display where the data can be viewed as a heatmap, scaled, and features can be added or removed either by the user or they can be filtered based on a threshold (see Fig. 4). For example, after scaling features, you could establish a threshold where features with a scaled value above 0.8 are removed from the analysis. Alternatively, you might be concerned about the number of NaN ("Not a Number") values being pruned or imputed in later analysis and so you could filter out the features with more than 25% missing values.

Describe data Enter a description of the data.

Figure 4: You can inspect the data in the matrix inspector to understand the data, identify errors, and exclude features or subjects. The descriptive statistics that can be displayed by using the drop-down list at the top right are the mean, median, variance, interquartile range, %(Nan) (i.e., "Not a Number" missing values), kurtosis, or skewness. The data can also be scaled using the button underneath the drop-down list. Once a descriptive option is selected, you can then enter a filter to exclude either F(Feats) or F(Cases) based on a value. When a feature filter is used, a red line will appear that signifies the filter value and features above this value will be excluded and moved to the "Feature pool" at the bottom left of the screen. You can manually add or subtract features or subjects using the arrows > between the feature pool and the features to import. Once features have been understood, filtered, or removed, press OK and return to the main menu.

IMPORT data This option will import data to the NeuroMiner structure ("NM") and progress to the next screen.

4.1.5 Matrix data: from a .mat file

This is the same as entering data from the workspace, but you can just select a pre-existing .mat file with all the variables. These are displayed at the top of the screen and then you can enter them using the directions above.

4.1.6 Matrix data: from a text file

The data can be stored in a .dat, .txt, or .csv file that has been saved with another program, such as Excel, R, or a database program. The file **must** contain labels for each of the IDs, labels, and features in the first row.

Define machine learning framework [classification]
Select data input format [matrix]
Define data provenance [Matrix from text file:]
Define delimiter in text file [tab]
Specify column header containing the label data [?]
Define column header containing the case IDs [?]
Inspect matrix data and select features for import
Describe data [?]
IMPORT

Define delimiter in text file Define the delimiter that was used to generate the file (i.e., comma, tab, etc.). It is critical to use the correct delimiter.

Specify column header containing the label data Write the column header variable name for the label data in your CSV file; e.g., "Illness" or "PANSS-General". For classification, the variable **must** be string variables instead of numerical coding (i.e., must be 'CTRL' and 'SCZ' instead of numbers representing these groups).

Define column header containing the case IDs Write name of the column header variable name for the case IDs in your CSV file; e.g., "ID". These must

be string variables (i.e., 'ID123' instead of 123).

Inspect matrix data and select features for import This feature opens a results display where the data can be viewed as a heatmap, scaled, and features can be added or removed either by the user or they can be filtered based on a threshold (see Fig. 4). For example, after scaling features, you could establish a threshold where features with a scaled value above 0.8 are removed from the analysis or features that have more than 10% NaN ("Not a Number") features are removed.

Describe data Enter a description of the data. This is important if you are entering multiple modalities.

IMPORT data Import data into the NM structure.

4.1.7 Matrix data: from a spreadsheet

This option will import data from a spreadsheet (only xls andxlsx). As with CSV files, the file **must** contain labels for each of the IDs, labels, and features in the first row.

```
Define machine learning framework [ classification ]
Select data input format [ matrix ]
Define data provenience [ Matrix from spreadsheet file: ]
Define name of sheet in spreadsheet file [ Tabelle1 ]
Specify column header containing the label data [ ? ]
Define column header containing the case IDs [ ? ]
Inspect matrix data and select features for import
Describe data [ ? ]
IMPORT
```

Define name of sheet in spreadsheet file Choose the sheet that you want to import.

Specify column header containing the label data Write the column header variable name for the label data in your spreadsheet file; e.g., "Illness" or "PANSS-General". For classification, the variable **must** be string variables instead of numerical coding (i.e., must be 'CTRL' and 'SCZ' instead of numbers representing these groups).

Define column header containing the case IDs Write name of the column header variable name for the case IDs in your spreadsheet file; e.g., "ID". These must be string variables (i.e., 'ID123' instead of 123).

Inspect matrix data and select features for import This feature opens a results display where the data can be viewed as a heatmap, scaled, and features can be added or removed either by the user or they can be filtered based on a threshold (see Fig. 4). For example, after scaling features, you could establish a threshold where features with a scaled value above 0.8 are removed from the analysis.

Describe data Enter a description of the data. This is important if you are entering multiple modalities.

IMPORT data Import data into the NM structure.

4.1.8 Add a Modality, Add Covariates, and Finish Analysis

Once data has been imported, you will be taken to the NM data workspace manager. Above the menu, you'll see a section called "MODALITIES" that will display the modality that has been entered and some descriptive statistics (e.g., the dimensionality and the number of NaN ('Not a Number') missing features). The menu will display the following:

- 1 : Add modality to NM workspace
- 2 : Modify modality in NM workspace
- 3 : Delete modality from NM workspace
- 4 : Delete all data from NM workspace
- 5 : Add covariate(s) to NM workspace
- 6: Delete covariate(s) from NM workspace
- 6 : Finish data import for discovery & cross-validation analysis

1: Add modality to NM workspace This option gives you the ability to add another dataset to the workspace in order for later data fusion or stacked generalization. It can be another modality (e.g., matrix data and then neuroimaging data) or it can be the same modality (e.g., structural neuroimaging data of one type and then of another type). The user will be then presented with the options that have been described above to repeat the process of data entry for another modality. It is critical to note that the subject identifiers have to be the same for each modality. In the case of NIFTI or SPM import, NeuroMiner will identify the subject IDs from the unique data in the file paths and thus any other data needs to have the same subject IDs (e.g., matrix data).

2: Modify modality in NM workspace This option takes you back to the "Input data into NM screen" where you can modify the settings or data described above.

3: Delete modality from NM workspace Delete a modality that has been entered.

4: Delete all data from NM workspace Delete all data and start again.

5: Add covariate(s) to NM workspace This is a critical feature of NeuroMiner that allows the entry of matrix data that can be used **in later preprocessing steps** to control for nuisance covariance (e.g., age and sex). It is **critical** to note that this step **simply enters the data and does not actually regress out the effects of the variables from your data**. However, these data can be used to correct data for nuisance covariates later in the preprocessing steps (see section [4.2.3.3](#)).

Please also note that if you want to normalise the data to a specific group

(see section 4.2.3.7), a binary vector (i.e., zeros and ones) must be entered at this stage containing ones for the individuals in the group that you want to normalise to.

→ **Critical Point**

The covariates must be stored in the MATLAB workspace—i.e., they must be loaded. They also must be a standard matrix double containing subjects in rows and covariates in columns. It is critical to note that the subjects **must be in the same order as the other data that has been entered**, regardless of what that data was. This is especially important if you are using an automatic method of selecting files for NIFTI images.

After the entry of the MATLAB matrix name, the variables can be defined one-by-one or by using a vector. The covariate names will then be displayed in the "MODALITIES" section at the top of the screen.

Delete covariate(s) from NM workspace Delete the covariates from the workspace.

Finish data import for discovery & cross-validation This option will finish the data import section. It is important to note that once this is selected **you will no longer be able to modify any part of the data**, including the addition of other modalities, covariates, or anything else. If there are changes, then you will need to enter the data again. Once the data has been processed, the user can however test models by entering out-of-sample data (e.g., external data from another center or group; see section).

4.2 Define parameter template

Once the data have been entered into NeuroMiner, the next step is to configure the settings that will be used to analyse the data. NeuroMiner does this by entering parameters (i.e., settings that define a system or sets the conditions of its operation) using the parameter template menu:

- 1: Data fusion [...]
- 2: Cross-validation settings [...]
- 3: Preprocessing pipeline [...]
- 4: Classification algorithm [...]
- 5: Learning algorithm parameters [...]
- 6: Ensemble generation strategies [...]
- 7: Visualization options [...]
- 8: Model saving options [...]
- 9: Save parameter template [...]
- 10: Load parameter template [...]
- 11: Define verbosity level [...]

In the case of machine learning analysis, the parameters are associated with all steps of the analysis and can be understood as a set of inter-dependent steps. These parameters are stored as a 'template' within the NM structure and analyses can be initiated using this template. Additionally, the template

can be saved (see 'Save parameter template') and imported to be used in other analyses.

It is important to note that all settings that are defined in the parameter template do not take effect on the data until they are 'initialized' using the main menu (see section 4). In other words, defining or changing a parameter template will not effect any existing analyses that have been initialized and/or run.

The main steps in a machine learning analysis can be characterised as: 1) defining the cross-validation parameters; 2) extracting features; 3) applying a classification algorithm (see section 4 and 3). Additionally, users may want to optimise the selection of variables using feature selection. These main steps are separately organised in NeuroMiner and will be outlined in the following subsections.

4.2.1 Data Fusion

Data fusion is the process of integrating different datasets or classifiers together (see [wiki](#)), ultimately with the hope that it will make your prediction better. For example, neuroimaging and clinical data could be fused together and a classifier could be used on these data. Alternatively, the researcher could conduct multiple analyses with different classifiers and then fuse the decision scores together. For theory of decision-based data fusion see the extensive work of Robi Polikar (e.g., [Polikar, 2006](#)) and for practical example see [Cabral et al., 2016](#). For data fusion to work in NeuroMiner, the user must enter different data sets that will be processed (called 'modalities' as discussed in section 4.1).

When different data modalities are entered, the first option in the parameter template will be: **1: Define data fusion options**. Selecting this option will lead to the following menu:

No fusion

Early fusion -> Modality concatenation BEFORE feature preprocessing

Intermediate fusion -> Modality concatenation AFTER preprocessing

Late fusion -> Decision-based data fusion

1 : No fusion Don't fuse the modalities.

2 : Early fusion -> Modality concatenation BEFORE feature preprocessing This is the most basic method of data fusion and means that the feature matrices will be concatenated and then all processing is conducted as normal. This will then display the modalities and ask the user to select the modalities that they would like to fuse. Modalities can be selected by entering single numbers (e.g., 1,2) or ranges (e.g., 1:3).

3 : Intermediate fusion -> Modality concatenation AFTER preprocessing The feature matrices are concatenated **after** they are **preprocessed** and then training is conducted. This will then display the modalities and ask the user to

select the modalities that they would like to fuse. Modalities can be selected by entering single numbers (e.g., 1,2) or ranges (e.g., 1:3). Keep in-mind that the free parameters defined in the separate processing of each modality will expand the overall model optimization process conducted in the training stage.

4 : Late fusion -> Decision-based data fusion Late fusion is when completely separate pipelines are run on each data modality, decision scores are produced, and then the scores are fused together and the decision score is based on an average of the decision scores. This is the most basic strategy for combining modalities at the classifier level and is known as bagging. This will then display the modalities and ask the user to select the modalities that they would like to fuse. Modalities can be selected by entering single numbers (e.g., 1,2) or ranges (e.g., 1:3).

Important note for intermediate and late fusion. If either intermediate or late fusion is selected, then it will be necessary to establish different settings (preprocessing, training, or both) for each of the different analyses. For this purpose, another menu item will be introduced into the parameter template called **2: Set active modality for configuration**. Selecting this option will display the modalities that are available. It is necessary to select one of these modalities, then to alter the preprocessing/training settings, and then move onto the next analysis. Following this process, you can then initialize the analyses and proceed with the analysis as normal.

It is important to note that fusing data will result in more complicated results that may not be as interpretable as a single analysis without fusion.

4.2.2 Cross-validation settings

Cross-validation is a fundamental core of machine learning analysis because it facilitates out-of-sample generalization and parameter tuning [link](#). NeuroMiner has been built to flexibly create custom cross-validation schemes depending on your data problem. NeuroMiner has been built around performing repeated, nested, cross-validation, which is a very robust method that increases the likelihood of generalization.

The cross-validation works by first defining settings of the following options:

- 1 : Select cross-validation framework [(Pooled) cross-validation]
- 2 : Define no. of CV2 permutations [P2 = 1]
- 3 : Define no. of CV2 folds [K2 = 10]
- 4 : Define no. of CV1 permutations [P1 = 1]
- 5 : Define no. of CV1 folds [K1 = 10]
- 6 : Equalize class sizes at the CV1 cycle by undersampling [no]
- 7 : Build CV2/CV1 structure
- 8 : Load CV structure
- 9 : Save CV structure

→ **Critical Point**

The user then must build the cross-validation structure in step 7 before exiting

the menu. The cross-validation framework will then be built within the NM structure.

4.2.2.1 Select cross-validation framework This option allows the user to select either k-fold cross-validation scheme (including leave-one-out) or a leave-group-out analysis with the following options:

- 1 — Pooled
- 2 — Outer Leave-Group-Out/Inner pooled
- 3 — Nested Leave-Group-Out
- 4 — Outer Leave-Group-Out/Inner Leave-Group-In

NeuroMiner has been built around the gold standard of repeated, nested cross-validation ([Filzmoser et al., 2009](#)) and therefore the default options are reflective of this. For a description of cross-validation, you could check-out the supplementary material of this publication: [Koutsouleris et al., 2016](#). In brief, leave-one-out cross-validation and k-fold cross-validation risk overfitting (i.e., your models not generalizing past your sample) in the context of (hyper)parameter optimisation (e.g., optimising C parameters) or the optimisation of feature selection based on the predictive accuracy on the held-out dataset (e.g., using wrappers discussed in the "ensemble" section of this manual). Nested cross-validation mitigates against overfitting and provides models that are more likely to generalise—which is the central aim of machine learning and science generally.

A depiction of nested cross-validation is represented in Figure 5. In NeuroMiner the outer cross-validation cycle of a nested cross-validation scheme is designated as "**CV2**" and the inner cross-validation folds are designated as "**CV1**". Models are trained in the CV1 cycle and then the best performing models are applied to the CV2 data. This separation of CV2 and CV1 data avoids overfitting.

Option "1 — Pooled" means that the outer and inner cross-validation folds will be automatically and randomly defined. Option "2 — Outer Leave-Group-Out/Inner pooled" means that the outer CV2 folds will be defined by the user using a vector defining groups (e.g., if there are different sites then this will test the ability of the classifiers that are trained with pooled site data to generalise to new sites). Option "3 — Nested Leave-Group-Out" means that both the outer and inner cycles will be separated based on the grouping vector that the user enters (e.g., in the site example, this will mean that the models will be optimised to generalise across sites in the CV1 cycle and then applied to different sites in the CV2 cycle). Option "4 — Outer Leave-Group-Out/Inner Leave-Group-In" means that the models in the inner CV1 cycle will only be trained in one group prior to being applied to all other groups (e.g., this will test the robustness of models from single sites). See 5 for more details and imagine that the folds are sites or diagnoses or any other grouping variable.

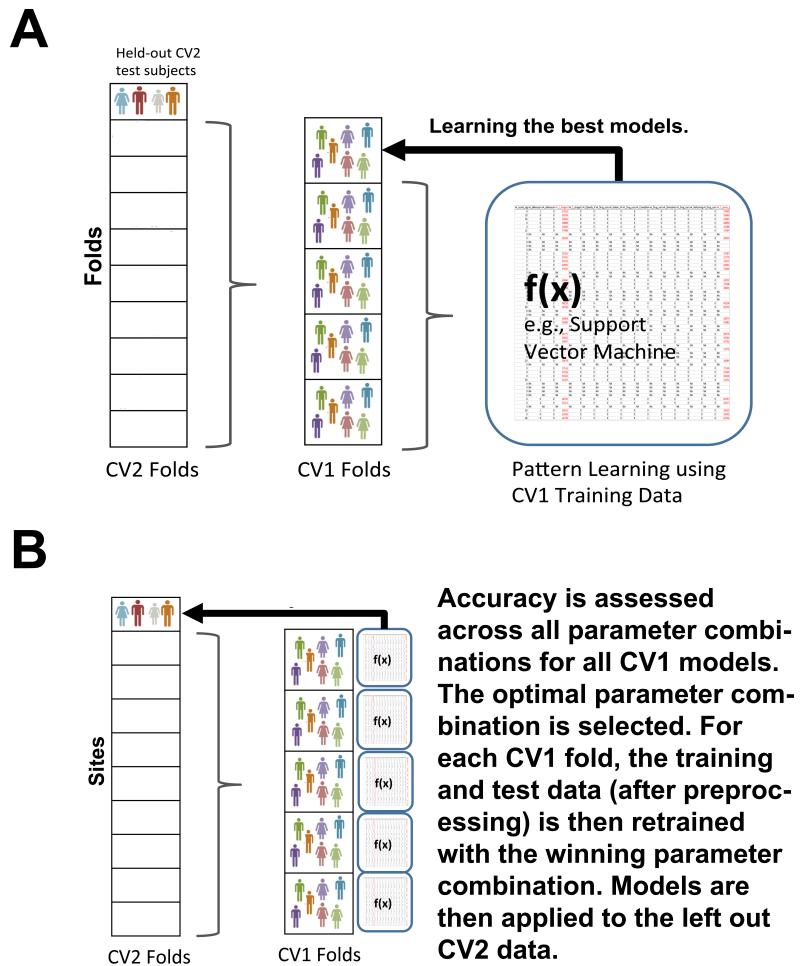


Figure 5: Nested cross-validation in NeuroMiner. **A:** The data is split into k number of folds in the outer CV2 set and a CV2 test set is held-out. The rest of the data goes into the 'nest' where it is again subdivided into k -folds, a test fold is held-out, and the rest of the data is used for training. Models associated with each parameter combination (e.g., C or ϵ) are created in the CV1 training data and applied to the CV1 test data. **B:** accuracy is assessed across all parameter combinations for all CV1 models. The optimal parameter combination is chosen based on a criteria set by the user. For each CV1 fold, the training and test data, which has been preprocessed, is retrained with the winning parameter or parameter combination. The models are then applied to the held-out CV2 test data. In cases there are many (hyper)parameter combinations or when there is variable selection procedure then the user can specify to keep a percentage of top performing models for each fold (this is discussed in the section of the manual on the learning algorithm parameters).

To further avoid overfitting and increase the generalizability of the model predictions, NeuroMiner also allows the user to implement ‘repeated’ cross-validation for each of the CV cycles. This involves shuffling the subjects prior to the definition of the folds so that there are different subjects in the test/training folds. This is usually done for both the CV2 and CV1 cycles.

A clear depiction of how repeated cross-validation works is difficult, but it is important to keep in-mind that there will be a lot of models produced under these schemes; e.g., if you have a standard 10x10 CV1/CV2 framework and you’re choosing the optimal model for each fold, then you’ll have 10,000 models. A way to understand how these are applied to an individual subject is to consider that all models in which that subject is **not** included in the CV1 training/testing are applied to them—because if they were included then this would be information leakage.

On the surface, having 10,000 models is not so great for the interpretation or understanding of a phenomena under investigation. However, doing this can be really good for generalization when it’s unlikely that there is a singular model that can be found and we need something that practically works—e.g., for many complex problems in psychiatry. For an interesting discussion on this topic, see [Breiman et al., 2001](#).

4.2.2.2 Define number of CV2 permutations Define the number of repeats that you want to happen to the outer, CV2 data. As stated above, the subjects are shuffled and new folds are created. The number of permutations is based on a balance between the computational time that your analysis is taking and the stability of the predictions from your modeling.

4.2.2.3 Define number of CV2 folds Define the number of folds that you want the outer data to be split into. If -1 is entered then a leave-one-out cross-validation is implemented and the permutations option is not possible (because it doesn’t matter if you shuffle the data if every single subject is left out). The number of folds is usually based on the number of subjects that you have.

4.2.2.4 Define number of CV1 permutations Define the number of inner, CV1 permutations in the inner folds. This is based on the computational time and also the model stability.

4.2.2.5 Define number of CV1 folds Define the number of inner, CV1 folds. The number of folds is usually based on the number of subjects you have and what kind of analyses that you are conducting. For an analytic example of where this is important see the section on ensemble generation strategies—e.g., using wrappers that are optimising features based on the subjects in the CV1 test folds.

4.2.2.6 Equalize class sizes at the CV1 cycle by undersampling Situations where there are unbalanced groups (e.g., less people transitioning to illness than people who do not) pose a problem for support vector machines and other analysis techniques. Usually, a model is learned that prefers the majority group and then you'll get unbalanced predictions (e.g., your specificity will be high, but sensitivity low). In NeuroMiner there are two ways to deal with this: 1) undersampling the majority group during training; 2) weighting the SVM hyperplane. This option allows the user to undersample the majority group so that the groups are balanced when the models are created in the CV1 folds.

After selecting the option, the user will be prompted with this: "Target ratio bigger / smaller class (1 - classes have sample size after balancing) (Default: 1.5)" This is the ratio of the bigger to smaller class sample sizes. Equal class sizes can be selected by entering in "1". The default is 1.5 because otherwise the models may not generalise well.

The next option gives the user the option to add the excluded CV1 subjects in the training folds to the CV1 test folds. This option is potentially important for potential generalizability of the models to the CV2 folds (and ultimately to real-life).

4.2.2.7 Build CV2/CV1 structure Once the options have been selected above, then the user needs to build the CV2/CV1 structure. This literally creates all the folds and indexes the subjects that are to be included in the folds. These can be found in the NM structure (e.g., NM.cv).

4.2.2.8 Load CV structure If you have saved the CV structure for your data with the subjects indexed into each fold/permuation, you can load it using this function. This can be useful if you establish a new analysis with new settings, but you want to load the CV structure that you made before.

4.2.2.9 Save CV structure This gives the user the option to save the indexed CV structure containing the folds and permutations for later use. It is saved in a .mat file.

4.2.3 Preprocessing pipeline

An important part of any machine learning analysis is how the data is prepared or 'preprocessed' prior to its analysis with a classification algorithm. This is also known as 'feature extraction' because the features are extracted from the existing data before analysis. NeuroMiner has a number of options to prepare data and can be tailored to a users specific data problem.

It's important to note that NeuroMiner performs preprocessing steps within the cross-validation framework. This means that when the option to preprocess the data is selected, it preprocesses the training data and applies the 'learned' preprocessing parameters to the CV1 and CV2 test data partitions.

When the preprocessing module is loaded the user will see the following:

```
=====
CV-PREPROCESSING SEQUENCE
=====
Step 1: Scale [ from 0 to 1 ], zero-out completely non-finite features
=> Step 2: Prune non-informative columns from matrix [ Zero Var, Nan,
Inf ]
You are here: MAIN
>>> Define parameter template >>>
-----
Preprocessing sequence generator: 2 / 2 steps.
-----
1 : Add preprocessing step
2 : Remove preprocessing step
3 : Insert preprocessing step
4 : Replace current preprocessing step
5 : Modify current preprocessing step
6 : << Previous step
7 : Change order of preprocessing steps
<== Back/Quit [Q]
```

Importantly, the preprocessing steps that are selected will be outlined in **in the order of analysis** in under the heading **CV-Preprocessing Sequence**. The user can add steps using the add preprocessing step and can modify steps using the other straightforward options within this menu—e.g., the user can modify the settings of a preprocessing step or can change the order of the steps.

4.2.3.1 Add preprocessing step This option is to add a preprocessing step to the "Preprocessing sequence generator". Once this is selected, the user can select the option that they would like from the following list:

- 1 : Regress out nuisance covariates
- 2 : Apply dimensionality reduction method to data
- 3 : Standardize data
- 4 : Scale data
- 5 : Normalize to group mean
- 6 : Normalize data to unit vector
- 7 : Apply binning method to data
- 8 : Impute missing values (120 NaNs in 25.0% of cases)
- 9 : Prune non-informative columns from data matrix
- 10 : Correct group offsets from global group mean
- 11 : Rank / Weight features
- 12 : Extract variance components from data
- 13 : Measure deviation from normative data

If you are using volumetric neuroimaging data, then there will also be an option included at the top of the list to: 1 : Enable spatial operations using Spatial OP Wizard. This is a special module designed to optimise filtering and smoothing within the cross-validation process and is always performed before any other preprocessing steps across the entire dataset (i.e., on test and training data).

Once the user has selected one of the steps, for example to perform a dimensionality reduction, they are then redirected to the main preprocessing menu that will add the processing step; for example:

```
=====
CV-PREPROCESSING SEQUENCE
=====
Step 1: Scale [ from 0 to 1 ], zero-out completely non-finite features
Step 2: Prune non-informative columns from matrix [ Zero Var, Nan, Inf ]
>> Step 3: Reduce dimensionality (PCA) [ dimensions: 0.8 ]
```

You can see that now there is a line stating CV-PREPROCESSING SEQUENCE containing "Step 3: Dimensionality reduction". The "Preprocessing sequence generator" indicates that the step is selected for further operations as indicated 3/3 (one out of a total of three) steps and this selection is also highlighted by the arrow symbols (>>). If an option requires a suboption (e.g., see section 4.2.3.4) then the suboption will be listed underneath the parent option and this can be selected using the arrows as well. As such, the user can now perform other operations within this menu, including removing the selected preprocessing step, inserting another preprocessing step at the same location, replacing the current preprocessing step with another one, or modifying the current preprocessing step.

If you have added a spatial preprocessing step, then this will appear above this menu because it is conducted prior to the preprocessing sequence across the entire dataset.

→ **Critical Point**

The preprocessing steps will be conducted in the order that they appear in the CV-PREPROCESSING SEQUENCE menu.

4.2.3.2 Enable spatial operations using Spatial OP Wizard You should observe the following:

1 — Select spatial operation [No filtering]

On turning on this mode, you will see the following menu:

- 1 : No filtering
- 2 : Absolute difference filtering (6 neighbors)
- 3 : Cube variance filtering (27 neighbors)
- 4 : Gaussian smoothing

Once an option is selected and the parameters have been defined, you will

see a new field above the "CV-PREPROCESSING SEQUENCE" in blue font as follows:

=====

NON-CV PREPROCESSING STEPS

=====

Resampling [3 Params: 3 (first) -> 9 (last)]

In NeuroMiner, it is designated as "NON-CV PREPROCESSING" because it occurs across the entire dataset before preprocessing. However, as described below, it's important to note that for smoothing the user can select a range of parameters and then the optimal combination of smoothing parameters, pre-processing settings, and training settings is found.

2: Absolute difference filtering (6 neighbors)

Absolute difference filtering is when the difference is computed between the voxel and each of the 6 nearest neighbors surrounding it. Then the value of the voxel is divided by the summed differences of the neighbors.

3: Cube variance filtering (27 neighbors)

The variance of the 27 neighbors surrounding each voxel is calculated and then the intensity of the target voxel is multiplied by the inverse of this variance.

4: Gaussian smoothing (=>FWHM)

This is regular Gaussian smoothing as used in most neuroimaging toolboxes. The advantage of doing this in NeuroMiner is that you can specify multiple different smoothing kernels (e.g., [6 8 10]) and then these will be used as hyperparameters during optimisation. That is, during learning in CV1 folds, the best combination of smoothing, other preprocessing steps, and learning parameters will be determined and applied to the held-out CV2 fold.

5: Resampling (=>Voxel size)

This is regular voxel resampling as used in neuroimaging toolboxes. In similarity to smoothing, the advantage of having it in NeuroMiner is that you can optimise across different resampling parameters to find the best combination of resampling, preprocessing, and training.

A critical point to remember in this mode is that if the user is interested to observe the Z-score by using the Visualisation option, it is essential for the user to use smooth data first. Applying a scaling model to unsmoothed imaging data will likely result in extremely high/low values in the test data due to non-overlap of tissue borders between training and test cases.

4.2.3.3 Regress out nuisance covariates If covariates have been entered when the data was entered into NeuroMiner (see [4.1.8](#)), then you can apply the correction using this step. This option is designed to remove the variance

associated with a nuisance variable (e.g., age, sex, study center) from the data within each CV fold. When chosen it will reveal the following options:

- 1 : Select covariates [age]
- 2 : Include intercept [yes]
- 3 : Attenuate or Increase covariate effects [attenuate]
- 4 : Use externally-computed beta coefficients [no]
- 5 : Define subgroup of training cases for computing betas [no]

Selecting the first option will allow you to choose the covariates that you want to control the data with:

(1) age

(2) sex

Select covariate(s) (enter expression - integer(s))

Select covariate(s) (Default: 1) :

You select the covariate(s) by entering in either a single numeral (e.g., 1) or many (e.g., 1:2) relating to the covariates that have been previously entered. Once these are selected, the user will be returned to the partial correlations setup menu.

2: include intercept. The choice of whether to include an intercept is based on your research question and relate to intercept inclusion in any other use of regression.

3: attenuate or increase covariate effects. Allows the user to either attenuate or increase the effects of the covariate on the data using regression.

4: Use externally-computed beta coefficients. Allows the user to enter beta coefficients from a regression that has been previously calculated. This option only works if the dimensionality of the beta coefficients exactly matches the dimensionality of the data. This means that NM will crash if the dimensionality of the data is dynamically changed during previous preprocessing steps.

5: Define subgroup of training cases for computing betas. Option 5 allows the user to define a subgroup of training cases for computing beta coefficients, which are then applied to the data (as discussed in the supplementary material of [Koutsouleris et al., 2015](#)). This function is useful when the user does not want to remove variance that may interact with the target of interest. For example, brain volume in schizophrenia interacts with age, therefore, the relationship between age and brain volume can be modeled in the control participants only and then these beta coefficients can be used in the schizophrenia sample to remove the effects of age without removing the effects of illness. When this option is selected, the user will see the following:

Compute beta(s) only from a specific subgroup (e.g. HC)?

yes / no (Default: no)?

Define index vector (ones = used / zeros = unused) for beta computation :

Here, the user needs to enter in a MATLAB logical vector consisting of TRUE

and FALSE corresponding to the participants in the study; e.g., [0 0 0 0 1 1 1]. A FALSE value (i.e., a zero) indicates that they will not be used in the calculation of the betas. A TRUE value (i.e., a 1) indicates that they will be used in the calculation of the betas. Logical vectors can be created from normal double vectors by simply typing "logical(yourvector)" on the command line.

4.2.3.4 Apply dimensionality reduction method to data A common need in machine learning analysis is to reduce the dimensionality of the data within the cross-validation framework (e.g., with structural neuroimages containing about 50,000 voxels). NeuroMiner allows the user to do this using a number of different methods.

When the option to reduce dimensionality is selected, the user will be shown the following menu:

- 1 : Principal Component Analysis (PCA)
- 2 : Robust Principal Component Analysis
- 3 : Non-negative Matrix Factorization
- 4 : Orthogonal Projective NMF (significantly slower but more robust compared to NMF) (currently visualisation mode is not integrated for the algorithm)
- 5 : Fast-than-fast NMF (fastest NMF through random projections and Nesterov iterations) (currently visualisation mode is not integrated for the algorithm)
- 6 : Partial Least Squares (PLS)
- 7 : Sparse PCA
- 8 : Probabilistic PCA (currently visualisation mode is not integrated for the algorithm)
- 9 : Factor analysis
- 10 : Locality Preserving Projections
- 11 : Linear Local Tangent Space Alignment (currently visualisation mode is not integrated for the algorithm)
- 12 : Large-Margin Nearest Neighbour (currently visualisation mode is not integrated for the algorithm)
- 13 : Linear Discriminant Analysis (low-D only) (currently visualisation mode is not integrated for the algorithm)

- 13 : Neighborhood Component Analysis
- 15 : * General Discriminant Analysis (low-D only)
- 16 : * Fast Maximum Variance Unfolding (currently not functional)
- 17 : * Laplacian Eigenmaps (currently not functional)
- 18 : * Local Linear Embedding (currently not functional)

Each of these options will ask for parameters that are specific to the type of

dimensionality reduction being conducted and are included in the respective sites for each technique listed below.

The main variable that can be changed across dimensionality reduction types is the number of dimensions that are retained following reduction (e.g., retaining 10 PCA components following dimensionality reduction of neuroimaging data). The following example applies to PCA reduction because this is the most common and we have found that it produces robust results.

For PCA, NeuroMiner gives the option to select the number of dimensions with the following options:

- 1 : Define extraction mode for PCA [Energy range]
- 2 : Define extraction range [0.8]

These options allow the user to select how the components are extracted and an extraction range based on this setting. When option 1 is selected, you will see the following menu:

- 1 : Absolute number range [1 ... n] of eigenvectors
- 2 : Percentage range [0 ... 1] of max dimensionality
- 3 : Energy range [0 ... 1] of maximum decomposition

Absolute number range [1 ... n] of eigenvectors. A whole number of components that they would like to retain a priori.

Percentage range [0 ... 1] of max dimensionality. A percentage of components to keep out of the total number of components.

Energy range [0 ... 1] of maximum decomposition. Retaining components based on the percentage of the total amount of variance explained (i.e., energy). For example, you might want to keep all components that explain 80% of the variance of your data.

For each of these options, except PLS, NeuroMiner gives the user the option to optimise the number of components that are selected during cross-validation by specifying a range of values. For example, a range of 20%, 40%, and 60% can be selected by entering [0.2 0.4 0.6]. NeuroMiner will then conduct all optimization procedures using these percentages of retained variance—i.e., it will find the best PCA reduction considering the other settings. If a range of values is required, then it is recommended that an additional substep is performed.

Extracting Subspaces

NeuroMiner as an additional option to first conduct the decomposition and then to retain different component numbers for further analysis. You do this by following the above procedure, but specifying that you want to retain a singular value of components (e.g., for PCA, 100% of the energy is recommended for the next step). Then you return to the preprocessing menu and select the option to "Add a preprocessing step". In the list of steps, there will now be an

option to "Extract subspaces from reduced data projections" and the following menu will appear:

- 1 : Define extraction mode for PCA [Energy range]
- 2 : Define extraction range [1]

Using these functions you can then select the components that you want to retain without having to run separate dimensionality reduction analyses. For example, you can retain 100% of the energy in the first step of a PCA to reduce the data dimensions and then retain a range of subspaces of components in the second step; e.g., [0.2 0.4 0.6 0.8]. This dramatically increases processing time.

The following advanced dimensionality reduction techniques are also offered by NeuroMiner, and the settings will change based on the technique. We recommend to follow the links provided below to determine the settings that are required or to evaluate the default settings offered within NeuroMiner.

Dimensionality Reduction Techniques

1. For all techniques see nk_PerfRedObj
2. Principal Component Analysis (PCA) – [PCA by Deng Cai](#)
3. Robust Principal Component Analysis – [LIBRA](#)
Note that RPCA is significantly slower than PCA but more accurate in low N problems.
4. Non-negative Matrix Factorization – [Li and Ngom's NMF toolbox](#)
Note that NMF is significantly slower than PCA but more parsimonious & robust.
5. Partial Least Squares
performs a single value decomposition (SVD; matlab built-in) on a covariance matrix constructed from the entered features plus another feature set to get latent variables. For example, if the primary features are voxels from a brain and the secondary PLS feature set are the labels it conduct SVD on the combined matrix (after standardization). It then multiplies this unitary matrix (U) from the SVD with the original standardized primary feature matrix in order to sensitise the analysis to the combination of the features. Other data can be used in place of the labels. The user has the option to use sparse PLS based on the function "spls".
6. Sparse PCA – Please check [preproc/spca.m](#); Zou et al., 2004
This is a 'statistical segmentation' tool.
7. Simple Principal Component Analysis – [MTDR](#); van der Maaten
8. Probabilistic PCA – [MTDR](#); van der Maaten
This is for low dimensional problems only, it requires >32GB RAM for high-D data.

9. Factor analysis – [MTDR; van der Maaten](#)
 10. Locality Preserving Projections – [MTDR; van der Maaten](#)
This is for low-dimensional problems only. Number of cases needs to be > no. features.
 11. Linear Local Tangent Space Alignment – [MTDR; van der Maaten](#)
For low-D only, requires >32 GB RAM for high-D data.
 12. Large-Margin Nearest Neighbour – [MTDR; van der Maaten](#)
 13. Deep Autoencoder – [MTDR; van der Maaten](#)
Low dimensional data only.
 14. Neighborhood Component Analysis [MTDR; van der Maaten](#)
Low dimensional data only.
- * For implementation of all functions, please see nk_PerfRedObj.m and files in the preproc directory.

4.2.3.5 Standardize data The standardization, scaling, and/or normalisation of data is an important part of most machine learning. NeuroMiner allows contains different methods to perform these functions within the preprocessing modules, which are conducted within each inner cross-validation fold.

Standardization is conducted on each of the features that have been entered to NeuroMiner across observations.

When option 3 is selected, the user will see the following:

- 1 : Select standardization method
- 2 : Compute standardization using a subgroup of cases [no]
- 3 : Apply standardization model to a subgroup of cases [no]
- 4 : Winsorize data (clamping of outliers) [no]
- 5 : Zero-out completely non-finite features [yes]

1 : Select standardization method

Select from the following standardization options: 1 — standardization using the median
2 — standardization using the mean
3 — mean-centering
4 — l1-median centering
5 — qn-standardization
6 — sn-standardization

Option 1 calculates the Z-score using the median (i.e., (score -median) / standard deviation) and option 2 calculates the Z-score using the mean. Option 3 mean centers all features (i.e., subtraction of the mean for the feature across all scores). Option 4 l1-median centering calculates the multivariate L1-median and then takes the derivatives from this as features (see

nk_PerfStandardizeObj.m; Hossjer and Crouse (1995) Generalizing univariate signed rank statistics for testing and estimating a multivariate location parameter, Non-parametric statistics, 4, 293-308). Options 5 and 6 are alternatives to the median absolute deviation for standardizing multivariate data (see Rousseeuw and Crouse (1993). Alternatives to the median absolute deviation, J. Am. Statist. Assoc., 88).

2 : Compute standardization using a subgroup of cases [no]

User can compute the standardization using a subgroup of cases.

3 : Apply standardization model to a subgroup of cases [no]

User can apply the standardization model to a subgroup of cases.

4 : Winsorize data (clamping of outliers) [no]

The user also has the option to Winsorize or 'censor' the data. This technique was introduced to account for outliers. After standardization, the elements which are outside of a user-defined range (e.g., 4 standard deviations) are set to their closest percentile. For example, data above the 95th percentile is set to the value of the 95th percentile. The feature is then re-centered using the new censored values.

5 : Zero-out completely non-finite features [yes]

If there are entries in the user's feature data that are non-numeric (i.e., "NaN" or "not a number" elements), then these will not be considered during the calculation of the mean and standard deviation (NeuroMiner uses the nanmean function). However, after the data has been standardized, these values will be added back to the feature matrix. As such, NeuroMiner gives the option to change these to zero in step four. It is important to note that once the data has been standardized, zero values reflect the mean of the feature and are thus a form of mean imputation. If the user wants to use other imputation options, they can select 'no' to this option and then impute the data using the preprocessing imputation option.

4.2.3.6 Scale data Another option to put the data into the same space is to scale it between two values. The user will see the following menu when they select this option:

- 1: Scale across features or each feature independently [each feature independently]
- 2: Define scale range [0,1]
- 3: Zero-out completely non-finite features [yes]

Using option 1, scaling can be conducted for "each feature independently" or "across the entire matrix". If each feature independently is selected, then the data will be scaled within the desired range for each feature (i.e., each voxel or each questionnaire). If the entire matrix is selected then each value in the

matrix will be scaled according to the range of all the data (e.g., all voxels or all questionnaires).

The user then has the option to define the scale range between 0 and 1 or between -1 to 1. These two options are provided because some machine learning algorithms require the data to be scaled differently; e.g., the liblinear options require the data to be scaled between -1 to 1.

Non-numeric values are not taken into account during scaling and are added back to the matrix after scaling. The third option gives the user the ability to either change these values to zero following scaling by selecting "yes" for the third option to "zero-out completely non-finite features". It is critical to note that in this case the values will be considered to be the absolute minimum of the scale for future analyses (e.g., if IQ is scaled, then these individuals will have an IQ value of 0). For questionnaire data, it is recommended that this option is changed to "no", and instead imputation is performed after scaling for the NaN elements. Alternatively, for imaging data, it is recommended that NaN values are excluded by using a brainmask during data input.

4.2.3.7 Normalize to group mean This option is only available when a group has been entered as a covariate at the start of the analysis (i.e., dummy coded vector containing ones for the group and zeros for the other participants). The option will then normalise all subjects to the mean of the specified group.

4.2.3.8 Normalize to unit vector This option takes the mean of the data, then subtracts this from each of the scores within the data. It then calculates the norm (L1 or L2) of this result. It then mean-centers the original data and divides each score by the normalised scores.

4.2.3.9 Apply binning method to data Binning is a technique that can be used to control for variance in data ([wiki](#)). For example, a histogram is an example of data binning. A form of it is discretization of continuous features ([wiki](#)) and this is offered in NeuroMiner as the default option when binning is selected as follows:

- 1 : Select binning method [discretize]
- 2 : Define binning start value [0]
- 3 : Define binning stepping [0.5]
- 4 : Define binning stop value [4]

Discretization basically just forms bins based on the standard deviation. You first establish a vector using a start value, a stepping value, and a stop value called alpha (e.g., the default setting is [0 0.5 1 1.5 2 2.5 3 3.5 4]). It then discretizes each feature by the mean +/- alpha*std. For example, for the first value of alpha, it finds the feature values that are above and below the mean and gives them a new value (i.e., 1 or -1). Then in the next step it finds the

values that are above and below the mean +/- half of the standard deviation and gives them a new value (i.e., 2 or -2). In the next step it finds values that are above and below one standard deviation from the mean and gives them a new value (i.e., 3 or -3). And so on. In this way, bins are created with a width that is determined by the standard deviation. It is important to note that this function will perform Winsorization (i.e., clamping) because values the stop value of alpha (e.g., above 4 standard deviations above or below the mean) will be given the final value.

Manually determining the bins is a good way to discretize data if you have a firm idea about the problem, but if you do not then it is useful to find an optimal bin width for a feature by taking into account the information content of the data. This can be done in NeuroMiner by changing the "1: Select binning method" to "Unsupervised entropy-based symbolization". It tries to find an optimal bin width for a feature by taking into account the entropy. It will do this within the range that you specify in the following settings:

- 1 : Select binning method [symbolize]
- 2 : Define minimum bin count [3]
- 3 : Define maximum bin count [25]
- 4 : Define sequence length [4]
- 5 : Define width of data range in number of STDs from mean [3]

4.2.3.10 Impute missing values The machine learning algorithms included in NeuroMiner are not able to process missing values (represented in MATLAB as "Not a Number" (NaN) entries). This means that they have to be either removed using the "Prune non-informative columns" feature described below, or they need to be imputed using this module.

NeuroMiner performs imputation using either single-subject median replacement, feature-wise mean replacement, or multivariate distance-based nearest-neighbor median imputation. When this feature is selected it will reveal the following menu:

- 1 : Define imputation method [kNN imputation (EUCLIDEAN)]
- 2 : Select features for imputation [All features]
- 3 : Define number of nearest neighbors [7 nearest neighbors]

1: Define imputation method

You have to first define the imputation method using the following menu:

- 1 : Median of non-NaN values in given case
- 2 : Mean of non-NaN values in given feature
- 3 : MANHATTAN distance-based nearest-neighbor search
- 4 : EUCLIDEAN distance-based nearest-neighbor search
- 5 : SEUCLIDEAN distance-based nearest-neighbor search
- 6 : COSINE similarity-based nearest-neighbor search
- 7 : HAMMING distance-based nearest-neighbor search

8 : JACCARD distance-based nearest-neighbor search

9 : Nearest-neighbor imputation using hybrid method

For median or mean replacement, values are imputed based on the data either across features for one subject or across subjects within the feature. For multivariate nearest-neighbor imputation, for each case with NaN values, a multivariate statistical technique is conducted to identify a number of similar (i.e., nearest-neighbor) cases from all the subjects that are available. For example, the Euclidean distance could be used to identify the 7 nearest-neighbor subjects that are close to the subject with the NaN value. Once these similar cases are identified in the multivariate space, then NeuroMiner will take the median of their values of the feature with the missing NaN case and impute it into the NaN field. A number of distance metrics can be used to represent the cases in multivariate space and to determine the nearest-neighbors based on the type of data you have (e.g., Manhattan, Euclidean, Seuclidean, Cosine, Hamming, or Jaccard). For example, the Euclidean distance could be used for data with a continuous measurement scale or the Hamming could be used for binary nominal data (i.e., 0 or 1) data (see Box: Nominal Data). There is also the option to use a hybrid approach that accounts for both continuous/ordinal and nominal data. It's important to note that all options except for the Manhattan and Euclidean distances require the Statistics and Machine Learning Toolbox of MATLAB.

1: Median of non-NaN values in given case

This function checks NaN values for each subject. When it finds a NaN value for a subject, it imputes the median of the non-NaN values across the other features for that subject. Therefore, this would make no sense at all if it was done across features that weren't equivalent in scale or from the same scale (e.g., a questionnaire). As such, it could be used in combination with the option to "2: Select features for imputation" described below.

2: Mean of non-NaN values in given feature

This function cycles through each feature and if it finds a NaN value then it imputes the mean of non-NaN values.

3: MANHATTAN distance-based nearest-neighbor search

This function determines the Manhattan distance between cases and then selects nearest-neighbors. It uses the 'distance' function from the Large Margins Nearest-Neighbors (LMNN) toolbox version 2.5 ([link](#)). You must scale, unit-normalize or standardize the data first, otherwise the distance measure will be dominated by high-variance features.

4: EUCLIDEAN distance-based nearest-neighbor search

Determines the Euclidean distance between cases and then selects nearest-neighbors. It uses the 'distance2' function from the Large Margins Nearest-Neighbors (LMNN) toolbox version 2.5 ([link](#)). You must scale, unit-normalize or standardize the data first, otherwise the distance measure will be dominated by high-variance features.

5: SEUCLIDEAN distance-based nearest-neighbor search

The Seuclidean distance between cases is measured using the pdist2 function that is built into the Statistics and Machine Learning Toolbox of MATLAB.

6 : COSINE similarity-based nearest-neighbor search

The Cosine distance between cases is measured using the pdist2 function that is built into the Statistics and Machine Learning Toolbox of MATLAB.

7 : HAMMING distance-based nearest-neighbor search

The Hamming distance between cases can be used for nominal data (e.g., 0/1 data) and is measured using the pdist2 function that is built into the Statistics and Machine Learning Toolbox of MATLAB.

8 : JACCARD distance-based nearest-neighbor search

The Jaccard distance between cases is measured using the pdist2 function that is built into the Statistics and Machine Learning Toolbox of MATLAB.

9 : Nearest-neighbor imputation using hybrid method

This option combines techniques that can be used for binary nominal data (i.e., 0/1) and ordinal/continuous data. You must define a method for nominal features (e.g., Hamming or Jaccard) and then define a method for ordinal/continuous features (e.g., Euclidean or Cosine). The maximum number of unique values for nominal feature imputation must also be inputted.

Box: Nominal Data. It is important to remember that nominal data with more than two categories needs to be dummy-coded for the results to make sense (i.e., 0 or 1 coding). NeuroMiner does not recognise the data-types and therefore nominal data that is coded with more than three categories (e.g., 1=Depression; 2=Bipolar; 3=Schizophrenia) will be used as a variable with meaningful relationships between the numbers when the differences are actually arbitrary.

2: Select features for imputation

The imputation procedures can be applied to the entire dataset or to subsets of the data. This is useful when there are different data domains (e.g., clinical and cognitive data) or even single questionnaires (e.g., the PANSS) that you want to restrict the imputation to. A MATLAB logical vector (i.e., consisting of TRUE and FALSE fields that are depicted as [0 0 0 1 1 0 0 etc.]; you can convert a binary vector to a logical simply by typing logical(yourvector) on the command line) must be provided to select a subset. If you want to impute multiple subsets then you must add separate imputation routines.

3: Define number of nearest-neighbors

If you're using the multivariate nearest-neighbors approach then you must define the number of nearest neighbors (e.g., subjects) to use for the calculation of the median value. The default is 7.

4.2.3.11 Prune non-informative columns from data matrix

Data pruning is the removal of examples or features that will not work with machine learning

algorithms or are redundant. In NeuroMiner, the machine learning algorithms that are used cannot work with missing values (i.e., "Not a Number" (NaN) entries in MATLAB) or with infinite values (i.e., "Inf" in MATLAB), which can occur due to previous processing steps. If the values of a feature do not change between participants (i.e., there is zero variance), then it is redundant to include the feature in the analysis as well. To correct for these problems, NeuroMiner offers the following options:

- 1 : Prune zero features [no]
- 2 : Prune features with NaNs [yes]
- 3 : Prune features with Insets [yes]
- 4 : Prune features with single-value percentage over cutoff [yes, 5]

The first option gives the user the ability to remove variables with no variance within a fold (i.e., if all individuals, or 'examples' within a cross-validation fold have exactly the same value for a feature).

The second option will remove features within a fold that have any NaN ('Not a Number') values. Similarly, the third option will allow users to remove features with any Insets (Infinity) values.

→ **Be careful with pruning**

It is important to note that even if there is one subject with one NaN or Inf within the fold NeuroMiner will remove the feature. The amount of removed features will be difficult to quantify. For matrix data, check the NaNs and potentially filter them prior to the entry into NeuroMiner (see [4](#)).

The fourth option is used in situations where there are a large number of subjects with the same value and a few with different values. For example, in neuroimaging due to registration inaccuracies most of the subjects may have a voxel value of zero in a location close to the gray matter border and a few might have non-zero numbers. In these circumstances, you might want to exclude the voxels where there is no variance in, for example, 90% of the sample (e.g., where 90% of individuals have a zero value). You can specify a percentage here to do this.

4.2.3.12 Rank / Weight features Selecting or ranking features is an important part of machine learning and can be achieved with **filters** and **wrappers**. For a discussion of why these methods are important and for all definitions see [Guyon et al., 2003](#). Wrappers will be discussed in section [4.2.6](#) of the manual, whereas filters are outlined here and also in section [4.2.6](#).

Filter methods broadly involve the weighting and ranking of variables, and optionally the selection of top performing variables. This can be conducted either prior to the evaluation of the model during pre-processing or as part of model optimisation (discussed in section [4.2.6](#)). This option allows the user to weight variables in the training folds based on some criteria (e.g., Pearson's correlation coefficient with the target variable) and then rank them in a specified order (e.g., descending order from most to least correlated). When this option is

selected, the following will be displayed:

- 1: Define target labels [NM target label]
- 2: Choose algorithm and specify its parameters [Pearson's]
- 3: Up- or down-weight predictive features [upweight features]

1: Define target labels

Here the user defines the target label that the filter will be applied to using the original data; e.g., this could be your group labels if you have a classification problem or your questionnaire item if you're using regression. The default setting in NeuroMiner is to use your target variable, but this can be changed by first selecting the menu item and then entering "D" to define a new target label. A new menu will appear and you can select to enter either categorical or continuous data. You can then enter a MATLAB vector that defines either categorical labels (e.g., [0 0 0 1 1 1]) or continuous labels (e.g., [15 26 19 33 22]). These variables should be stored in your MATLAB workspace.

When a continuous variable is selected, the user will then be asked whether they want to weight feature using only one specific subgroup. This means that the feature will only be weighted only based on the relationship between the feature and the filter target in one subgroup (e.g., control participants).

2: Choose algorithm and specify its parameters

The user here has the option to choose from among a number of methods that will weight the relationship between the filter target and the feature. These are:

- 1: Pearson / Spearman correlation
Standard MATLAB functions for Pearson's (corrcoef).
- 2: Spearman correlation
Standard MATLAB function for Spearman's (corr).
- 3: G-flip (Greedy Feature Flip)
See [Gilad-Bachrach et al. \(2004\)](#) and [site](#)
- 4: Simba
See [Gilad-Bachrach et al. \(2004\)](#) and [site](#)
- 5: IMRelief
This is an implementation of the Yijun Sun IMRelief function called "IMRelief Sigmoid FastImplementation" (see [Sun & Li, 2007](#)).
- 6: Relief for classification
This is a version of the built-in MATLAB ReliefF algorithm ([link](#)).
- 7: Linear SVC (LIBSVM)
This weights the features by conducting a linear SVM using each independent feature, as implemented in LIBSVM ([Chang & Lin](#)).
- 8: Linear SVC (LIBLINEAR)
Weights the features based on using LIBLINEAR toolbox ([link](#))

9: F-Score

Weights the features based on conducting an ANOVA using standard MATLAB functions on the differences between the groups when problem is classification.

10: AUC

Calculates an area under the curve (AUC) for binary classifications based on the code of Chih-Jen Lin and then ranks the features based on this.

11: FEAST

Implementation of the FEAST toolbox for feature selection ([site](#)).

12: ANOVA

Standard implementation of ANOVA.

13: External ranking

This option uses a user-defined weighting template that was entered during data input (see [4.1](#)), to use an external weighting map from a file, or to read-in a weighting vector using the MATLAB workspace.

3: Up- or down-weight predictive features

This option allows the user to either upweight or downweight the features based on the relationship between the variable and the target. For example, if you want to increase the importance of the features based on the target you would upweight the features. If you want to reduce the importance of the features based on some other variable (e.g., age if it is a nuisance variable) you could select the variable as the feature selection target and then choose the downweight option. For more information about why you would want to do this see references above.

4.2.3.13 Correct group offsets from global group mean This function was introduced primarily to control for scanner differences in structural neuroimaging data. In the case of two centers, for each feature, the mean of both center A and center B is calculated. The mean difference between these is then calculated ($A-B=Y$) and then subtracted from each of the data points in A (i.e., $A_i - Y$). This value is stored for further analysis. Internal empirical tests indicate that this basic method can control for major site differences.

4.2.3.14 Extract variance components from data This function takes a source matrix and performs a PCA. It then determines correlations between the eigenvariates from the PCA reduction, and those above or below a specified threshold are identified. It then projects the target matrix to the source matrix PCA space, and then back-reconstructs this to the original input space without the identified PCA components (i.e., it removes the variance associated with those components).

4.2.3.15 Measure deviation from normative data This function calculates a PLS and then determines how much the data deviates from the model that

is produced. The hypothesis here is that the PLS model will represent a normative variation between the primary features (e.g., the brain) and the second feature matrix (e.g., the labels or another variable set). By calculating the deviation from this model, it is hypothesised that normative deviation will be modeled and then included in further processing steps.

4.2.4 Classification algorithm

Support Vector methodologies are the base of the evaluation core of NeuroMiner, other methods such as k-Nearest Neighbors and Random Forests are also available. In this distribution we rely on the LIBSVM /LIBLINEAR toolboxes for the implementation of support vector based classifiers and regressors. matLearn is also now introduced and requires the Statistics and Optimization Toolbox.

NeuroMiner interface will change according to the evaluation type selected in the data input process, two operation modes are implemented: Classification and Regression. The following will refer to each of these separately.

4.2.4.1 Classification mode In this section we will describe the set up of **Classifiers** in Neurominer.

4.2.4.1.1 Define prediction performance criterion During machine learning, models will be optimised and evaluated based on a criterion. For classification problems NeuroMiner provides the following performances measures:

1. Accuracy: $\text{sum}(E = P) * 100 / N$ **ACC**
2. True Positive Rate: e.g. for one-class SVM
3. False Positive Rate: e.g. for one-class SVM
4. Positive Predictive Value: e.g. for one class SVM
5. Matthews Correlation Coefficient:
$$(TP * TN - FP * FN) / \sqrt{(TP + FP)(TP + FN)(TN + FP)(TN + FN)}$$
MCC
6. Area-Under-the-Curve **ROC**
7. Geometric Mean: $\sqrt{PPV * True\ Positive\ Rate}$
8. Balanced Accuracy: $(sensitivity + specificity) / 2$ **BAC**
9. Enhanced Balanced Accuracy: $(sensitivity * specificity) / 100$
10. F-score: $(1 + \beta^2) * precision * recall / (\beta^2 * (precision + recall)) * 100$
11. Prognostic Summary Index: $PPV + NPV - 100$
12. Number Needed to Predict: $1 / PPV$

Notes: E, expected; P, predicted; TP, true positive; TN, true negative; FP, false positive; FN, false negative; PPV, positive predictive value; NPV, negative predictive value.

Each of these performance measures can have an effect on the resulting model and is dependent on the analysis that is being conducted. For example, simple accuracy is useful when it doesn't matter if sensitivity and specificity needs to be balanced (e.g., if sensitivity could be very high and specificity very low). However, if there are unequal sample sizes then this can lead to very unbalanced sensitivity and specificity that is an artifact of sample size, which can be corrected by using balanced accuracy (BAC).

4.2.4.1.2 Configure learning algorithm Learning algorithm to be used for the **classification** analysis. Each algorithm has a particular set of parameters that are described in the respective distributions. In general defaults are provided.

1. SVM – [LIBSVM](#)
2. SVM/LR – [LIBLINEAR](#)
3. RVM / BAYES – [Mike Tipping's implementation](#)
4. Local Learning – [Akbas fuzzy k-Nearest Neighbors \(fKNN\)](#)
5. Univariate – GLM Logistic Regression [MATLAB glmfit](#)
6. Decision Tree – [MATLAB Statistics Toolbox implementation of the Decision Tree algorithm](#)
7. Random Forest – [Abhishek Jaiantilal's fast Random Forest algorithm for MATLAB](#)
8. matLearn – [matLearn](#). This only works when users have the Statistics and Optimization Toolbox
9. GLMNET – [Hastie's library for LASSO/Elastic-net regularized GLM](#)
10. Gradient Boost – [Carlos Becker's gradient boosting algorithm](#)

Once one of these option is selected, you will have to define a number of other settings that are specific to the algorithm and can be found in the respective references above. In the following, we present the options that are available for SVM.

SVM options

The following options will appear when SVM is selected:

- 1 : LIBSVM version [3.12]
- 2 : Classifier type [C-SVC (L1-regul.)]
- 3 : Probability estimation using Platt's method [no]
- 4 : Cache size [500 MB]

- 5 : Termination criterion [0.001]
- 6 : Shrinking heuristics [yes]
- 7 : Weighting of hyperplane in uneven group sizes [no]
- 8 : Case-level weighting vector [Case weighting not activated]

→ [Weight the hyperplane for uneven group sizes](#)

Each of the first 6 options are standard options within the LIBSVM toolbox [LIB-SVM](#). However, NeuroMiner also offers hyperplane weighting in uneven group sizes. This is an important option because support vector machines work effectively in circumstances with balanced samples of class examples but sub-optimal results can emerge when there are unequal class proportions; i.e., the hyperplane can be biased towards a correct classification of the majority class, thus resulting in poor classification performance in the minority class. This is called an imbalanced learning problem. To account for this problem class-weighted SVM can be used. Rather than having equal misclassification penalties (C values) for all individuals this approach corrects for imbalanced learning by increasing the misclassification penalty in the minority class. Specifically each C parameter for each class is multiplied by the inverse ratio of the training class sizes. To optimise the search for optimal class weights, you can also introduce an additional hyperparameter that multiplies the inverse ratio by a scaling factor for each class prior to the multiplication of the C parameter. This can be defined when the learning algorithm parameters settings in section [4.2.5](#). Similarly, option 8 allows you to weight the hyperplane based on a vector of weighting parameters that you have defined. This will work similarly to the above, but it will use the parameters entered rather than the group sizes.

Once these options are selected, NeuroMiner will progress to the next menu to define the kernel type.

4.2.4.1.3 Define kernel type Define whether the algorithm will use a linear or non-linear kernel [Kernel](#). This will display the following options:

- 1: Linear
- 2: Polynomial
- 3: RBF (Gaussian)
- 4: Sigmoid

Once these settings are defined, then you will be directed back to the main menu for this section and you can proceed to the main menu and then can configure the learning algorithm in section [4.2.5](#).

4.2.4.2 Regression mode If you have defined a regression framework during data input (see section [4.1](#)) then a different menu will appear as follows.

4.2.4.2.1 Define prediction performance criterion Performance criteria play a major roles in model evaluation and optimization, for regression NM provides 5 performances measures:

1. Mean Squared Error [MSE](#)
2. Normalized Root of Mean Squared Deviation [NRMSE](#)
3. Squared Correlation Coefficient [R²](#)
4. Correlation Coefficient [CC](#)
5. Mean Average Error [MAERR](#)

4.2.4.2.2 Select learning algorithm Learning algorithm to be used for the **regression** analysis. Each algorithm has a particular set of parameters that are described in the respective distributions. In general defaults are provided.

1. SVM – [LIBSVM](#)
2. SVM – [LIBLINEAR](#)
3. RVM / BAYES – [Mike Tipping's implementation](#)
4. Multinomial relevance vector regression (MVRVR) (for multiple target label prediction)
5. Univariate – GLM Linear Regression

4.2.4.2.3 Define kernel type Kernel to be used in the selected **regressor** (when applicable).

1. Linear
2. Polynomial
3. RBF (Gaussian)
4. Sigmoid

4.2.4.2.4 Optimize decision threshold based on ROC This option allows to optimise the decision scores according to the Receiver Operator Function. The default setting enables the post-hoc optimisation using ROC analysis of C1 test data.

4.2.4.2.5 Use ADASYN to adjust for unbalanced class settings Adaptive Synthetic (ADASYN) is primarily used to tackle the 'Imbalanced Classification Problem' prominent in Machine Learning which arises due to imbalanced class size. The key hindrance with an imbalanced dataset is that one class dominates another class resulting in a model that is highly underfitted as it is not able to classify the minority class successfully. ADASYN is an algorithm that generates synthetic data by generating more data for the class with lesser data.

Enabling ADASYN mode opens up extra options:

- 7: Define beta value (defines how much balancing will be applied, 0,1) [k=1]**
- 8: Define k value of density algorithm (kNN looking at both classes) [k=5]**
- 9: Define k value of SMOTE algorithm (kNN looking only at the minority class) [k=5]**
- 10: Data is already normalized for kNN search [no]**

4.2.4.3 Sequence Optimization When in the NeuroMiner "expert" mode (i.e., loading NeuroMiner with the command "nm nosplash expert") another option in the "model algorithm" menu is to conduct sequence optimization. This option was designed to simulate the clinical decision making associated with additional testing. In a clinical environment, the decision to conduct another assessment on an individual is normally done in situations of uncertainty to confirm a diagnosis—e.g., coronavirus is suspected based on a clinical assessment and then a nasal swab is conducted to confirm this diagnosis. In a machine learning context, uncertainty can be quantified with either the distance from the decision boundary in the case of an SVM (i.e., with cases closest to the decision line being the most uncertain) or with a median using any other scores.

NeuroMiner is able to test sequences of assessments (e.g., clinical, cognitive, imaging) based on the modalities entered during data entry or it is able to build optimal sequences based on different models (e.g., linear and non-linear SVM). The pipelines need to have been analysed before sequencing can be activated. The procedure is to:

1. Process multiple analyses as usual from different modalities or using different settings
2. In the "Parameter Template" menu, activate the option to 1: Define meta-learning/stacking options
3. Select the analyses that you want to define the optimal sequence for
4. Go to the preprocessing parameters and change the set-up so that there is only an imputation step. If you are sequentially stacking different learning algorithms, then scaling may be required and this needs to be considered
5. In the "Parameter Template" menu, go to the "Classification Algorithm" option and then the "Select learning algorithm" option
6. Select the item "SEQOPT —; Predictive sequence optimization algorithm (Stacking only)"
7. To configure the sequence algorithm, then select the option to "Configure Sequence Optimizer"

4.2.4.3.1 Optimization based on...

1. The population to be propagated
2. The entire population

NeuroMiner will attempt to optimise two factors: 1) the criterion that the user decides in the previous step (e.g., BAC); 2) the least number of assessments necessary. These optimizations can be for only the cases that have uncertain decision or probability scores (option 1) or for the entire population.

4.2.4.3.2 Define computation mode for combining sequential predictions

1. Replacement
2. Mean across predictions

When a case is propagated to another assessment because the first assessment was uncertain, then either their decision score can be replaced or a mean can be taken across the predictions in the sequence (i.e., this becomes bagging or "late fusion" in NeuroMiner terminology).

4.2.4.3.3 Define propagation algorithm's anchor

1. Decision boundary
2. Median

As mentioned above, either the decision boundary can be used to define uncertain cases (i.e., near the boundary) or the median of other scores can be taken. The first item uses the decision boundary and the second item uses the median to quantify the uncertainty (i.e., the median of scores that range from very likely group 1 to very likely group 2 is assumed to be uncertain here).

4.2.4.3.4 Define optimisation mode

1. Optimization criterion
2. Mean decision distance

Here the user can select whether to optimize on the criterion previously selected (e.g., the BAC) or on the mean decision distance (only works for SVM or algorithms with decision scores). In the second case, the idea is that the further away from the decision boundary the better the optimization.

4.2.4.3.5 Sequences to be tested

1. Define sequences matrix to be evaluated

Here the user can define what sequences are to be evaluated. This step is dependent on the identification of models to be included in the "1: Define meta-learning/stacking options". The idea is that rather than testing all combinations

of the models (e.g., model 1 - \downarrow model 2 - \downarrow model 3 and then all other combinations of these), the user can specify the models that they want to test. This is because usually there is a rationale for testing model sequences based on clinical experience and also cost (e.g., clinical assessment, then cognitive assessment, then imaging if necessary).

Sequences of assessments need to be specified here either directly or using a matlab variable. For example, if you wanted to specify a sequence of model 1 - \downarrow model 2 - \downarrow model 3 then you would enter [1 2 3]. If you wanted to specify this and also models 2 1 3 then you would enter [1 2 3; 2 1 3]. If you wanted to specify more then you would create a matlab variable with more sequences in rows and you would enter the variable name into this section.

4.2.5 Learning algorithm parameters

Parameters for a learning algorithm (e.g., an SVM or SVR) change the way that the model is fitted to the data. A common example is the "C" or "slack" parameter for soft margin SVM that defines the penalty associated with misclassifying individuals—i.e., whether the model does not allow any error within the CV1 folds and very tightly fits the decision boundary, or whether it allows error in order to improve generalizability. As such, parameters such as these can greatly affect the model performance and it is common practice in machine learning to optimise the parameters for your specific problem and data.

NeuroMiner was developed in order to optimise performance across pre-defined parameter ranges within the nested cross-validation framework using a gridsearch, which protects against overfitting due to the application of the trained models to held-out data. Parameter default ranges are provided in NeuroMiner based on the literature and empirical testing, but we strongly recommend that the parameters are defined for your study and problem.

As previously stated, NeuroMiner uses a dynamic menu configuration that changes based on previous input. This is no different for the learning algorithm parameters, whereby the menu options will change based on what you have selected in the "Classification algorithm" section.

Here we show an example for a RBF-Gaussian kernel Support Vector Machine **classifier**.

1. Define Slack/Regularization parameter(s)
2. Define Kernel parameter(s))
3. Define Weighting exponents
4. Enable regularization of model selection

Model selection regularization option controls the optimization module in NeuroMiner. When parameter(s) selection is necessary regularization might be

advisable. If you answer yes to this question the following options will be available.

1. Criterion for cross-parameter model selection
2. Suboption 1 dependent on (a)
3. Suboption 2 dependent on (a)
4. Specify cross-parameter model selection process. This option controls the number of models (parameters combination) selected. One optimal model based on the criteria and regularization defined previously or an ensemble of the top performing models

4.2.6 Ensemble generation strategies

NeuroMiner automatically works with ensemble learning at the CV2 level across the CV1 folds of a nested, repeated cross-validation design (see Fig. 5), but it also includes the ability to apply filters and wrappers that optimise predictions by selecting features within these folds (for more information, see [link](#) and [Kolhavi & John](#)).

4.2.6.1 Filter-based ensemble generation NeuroMiner has two methods of applying filters: 1) as part of the preprocessing of features by pre-selecting variables based on their relationship to the target (see 4.2.3.12); 2) as an ensemble method where different variable sets (variable subspaces) can be chosen based on the performance of the machine learning training algorithm. This latter case is known as constrained feature optimization because the variable sets are produced based on the performance of an algorithm (e.g., correlation), the sets are evaluated with the algorithm, and then the models corresponding to well-performing variable sets are chosen as an ensemble.

The first step is to turn the option to yes specifying the "train filter methods on CV1 partitions". This will generate a menu with the following:

- 1: Train filter methods on CV1 partitions
- 2: Use algorithm output scores or label predictions
- 3: Specify filter type
- 4: Define minimum number of features to be selected
- 5: Specify subspace optimization strategy
- 6: Define target population for computing optimization
- 7: Define subspace stepping

1: Train filter methods on CV1 partitions Turn constrained feature optimization filtering on or off.

2: Use algorithm output scores or label predictions Choose the feature sets based on majority voting of the predicted labels (hard decision ensemble) or based on an average of the predicted probabilities from the machine learning algorithm (soft decision ensemble).

3: Specify filter type Specify what filter will be used to obtain the feature sets. These methods are the same as in the preprocessing filter outlined in section [4.2.3.12](#).

4: Define minimum number of features to be selected This gives the option to define a minimum number of features that are retained when choosing feature sets.

5: Specify subspace optimization strategy Once the subspaces are defined and the models are evaluated, this section will define which subspaces/models are retained. There is the option to retain one model of the best performing subspace, or an ensemble of models across well-performing subspaces.

1: Subspace with maximum argmax. Argmax (arguments of the maxima) are the points of the domain of some function at which the function values are maximised. This option picks the maximum argmax across the feature subspaces (winner-takes-all).

2: Subspace ensemble with maximum argmax Picks the X most predictive subspaces with reference to the argmax.

3: Subspace ensemble with maximum argmax above a percentile Picks the top X% of subspaces.

4: All-subspace ensemble Uses all subspaces when training the model.

6: Define target population for computing optimization You can choose the best performing models using the CV1 training data, the CV1 test data, or the CV1 training & the test data. This means that the subspaces will be selected on the basis of how they predict the labels in each of these data folds. Selection based on training and test data is recommended.

7: Define subspace stepping This option defines how the subspaces are formed. The features are ranked based on the association between them and the target variable, and then they are divided into subspaces based on blocks of X% of features; e.g., blocks of 10% of features would divide the data into the top 10% performing features, then the top 20% of features, then the top 30% of features, and so on.

It is important to note that once the feature subspaces are defined, the features within the winning subspaces can then be used in a wrapper to further optimise performance.

4.2.6.2 Wrapper-based model selection The wrapper methods in NeuroMiner use either Greedy feature selection or simulated annealing to select feature combinations that maximise the predictive accuracy of a model in the CV1 data. You will see the following menu:

1 : Deactivate wrapper methods

2 : Wrapper type [Greedy feature selection (Use CV1 test data; Stop at k=50% of features; Cross-CV1 PFE off)]

3 : CV1 data partitions for optimization [CV1 test data]

4 : Cross-CV1 Feature selection [Cross-CV1 PFE off]

2: Wrapper type Select Greedy feature selection or simulated annealing.

Each option will lead to a new menu as follows:

Greedy Feature Selection

1 : Search direction [Forward]

2 : Early stopping [Stop at 50% of feature pool]

3 : Feature stepping [Each feature will be evaluated]

The search direction must be provided (forward or backward), which determines whether an empty subspace is filled with features that are the most predictive of the target variable (forward selection) or whether all the variables are entered into the subspace and then variables are removed until the optimal model is found (backward)

Early stopping is when variables are added up to a total percentage of the feature pool and then no more variables are added. This feature is useful when you want to find parsimonious solutions and also because restricting the feature pool can lead to better solutions.

Feature stepping can also be changed from adding/removing single features and then testing the model, or adding/removing percentages of features before testing the model. Adding percentages (e.g., 5%) normally results in a faster processing time.

Simulated Annealing

1 : Sparsity constant [0.01]

2 : Starting temperature [1]

3 : Stopping temperature [0.005]

4 : Alpha [0.9]

5 : Maximum # of iterations [2000]

6 : Maximum # of repetitions in Temperature T [100]

7 : Minimum # of repetitions until solution is accepted in T [30]

8 : k constant (smaller k ==> less solutions accepted) [10]

Simulated annealing is a probabilistic technique for approximating the global maximum of a function [wiki](#). The settings are standard parameters for any simulated annealing analysis.

3: CV1 data partitions for optimization Choose either the CV1 training, test, or test & training data to optimize the feature set.

4: Cross-CV1 feature selection In a nested cross-validation design, you will have a selection of features for each of the CV1 partitions. This step allows you to select features across the CV1 folds and will reveal another menu:

1 : Optimize feature selection across CV1 partitions [yes]

2 : Probabilistic feature extraction mode [% Cross-CV1 feature selection agreement]

3 : Apply consistency-based ranking to [Prune unselected features from each model and retrain models]

Probabilistic feature extraction

1 : % Cross-CV1 feature selection agreement

2 : Absolute number of most consistently selected features

3 : Percentage of most consistently selected features

The first option allows you to select features that occur across CV1 partitions at a certain percentage rate. For example, select only features that appear across CV1 folds 75% of the time or above this. Sometimes the threshold that is set does not return any features, and therefore you will also be asked to define a tolerance value for this circumstance. For example, if the 75% criterion is not met, then reduce this by 25% (i.e., so then you are effectively selecting features 50% of the time). You will also be asked to define a minimum number of features that must be selected each time (e.g., 1 feature).

Option 2, "Absolute number of most consistently selected features", just orders the features based on the amount of times they were selected across the CV1 folds. Then you can select however many features that occur at the top of the list (e.g., you could select the top 10 most consistently selected features).

Option 3, "Percentage of most consistently selected features". This option sorts in the same way as described for absolute number of consistently selected features and then establishes a cut-off. The features above this cut-off are kept. For example, if 90% is the cut-off then it will select the top 10% most consistently selected features across the CV1 folds.

Apply consistency-based ranking

1 : Retrain all CV1 models after pruning them from unselected features

2 : Retrain all CV1 models using the same selected feature space

Once the features are selected, you can either "1: Retrain all CV1 models" after pruning the unselected features. Or you can completely retrain the models using the single optimized feature set that was found in the previous steps.

4.2.7 Visualization options

In NeuroMiner, visualization refers to the representation of the model weights for each feature. In the case of neuroimaging, this means that the brain maps are 'visualized'. This setting gives the user the option to derive Z scores and P values using permutation analysis.

The magnitude of the feature weights do not necessarily inform the user about the statistical significance (i.e., the degree to which the result may replicate)(see [Bilwaj & Davatsikos, 2013](#)). To do this, permutations analysis can be conducted to create an empirical null distribution of weights for each feature and then the observed weight is compared to this distribution. This is described in the supplementary material of [Koutsouleris et al. \(2016\)](#).

The method is based on [Golland & Fischl \(2003\)](#) and involves the permutation of the outcome labels, features, or both. For each permutation, the models are retrained in the cross-validation framework using the respective feature/label subsets obtained from the observed-label analyses. For each permutation, the predictions are accumulated of the random models into a permuted ensemble prediction for each CV2 subject. Thus, a null distribution of out-of-training classification performance (BAC) for the prediction models is constructed. The significance of the observed out-of-training BAC is calculated as the number of events where the permuted out-of-training BAC is higher or equal to the observed BAC divided by the number of permutations performed. The significance of the model can then be determined by using a p threshold (e.g., $p>0.05$).

When the option is selected, you will be prompted to enter the number of permutations conducted (this will be dependent on sample size, analysis complexity, and computational time) and then whether you would like to have the labels, the features, or both permuted (see Golland et al., 2003 for details).

It is important to note that using PCA with this method will result in overly conservative p-values and is not recommended.

4.2.8 Model saving options

This option will save all the models used in training and testing. This is generally not a good idea because it will take up too much space.

4.2.9 Save parameter template

Once the parameters are defined in the above steps, they will be included within the "NM" MATLAB structure and will be ready to be "initialised" in the previous menu. However, the user also has the option here of saving the parameter settings into a separate .mat file so that they could re-load them for another analysis. This allows the user to be certain that the same parameters are applied in analyses with different data.

Please note when saving that loading parameter templates from different versions of NeuroMiner is currently unsupported.

4.2.10 Load training template

When a parameter template has been saved in the previous menu item, it can be loaded with this feature. This will parametise all settings within the template of the new NM structure, and thus will allow it to be initialised in the next step.

Please note that you will see a menu item warning about the cross-validation set-up. This warning is to highlight the fact that if you have different data then you need to re-create the cross-validation structure.

4.2.11 Define verbosity level

During preprocessing, this option will either output all processing steps being performed on the data using the "detailed" setting (recommended for first time use and to check processes) or you can set it so that it only outputs the training fold that it's working on and nothing else (the "none" setting). Setting this to "none" will increase the speed of analysis because it will not have to print the steps to screen.

4.3 Multi-group analyses

Multi-group analysis is a more advanced feature of NeuroMiner where classification is used to distinguish between subjects from multiple groups (e.g., schizophrenia, bipolar, and depression subjects). During data input, these groups are simply defined within the labels—e.g., 'SCZ', 'BP', and 'DEP' identifiers. Once this is defined, then an option will appear in the parameter template called "Multi-class settings". Selecting this option and then selecting the option to "Train multi-class predictor" to 'yes' will reveal another option to "3 : Specify multi-class decision mechanism []". This option determines how NeuroMiner deals with the multiple groups (e.g., how it optimises the predictions) and is critical for the interpretation of results. For all the methods included here, it is important to note that NeuroMiner works by conducting multiple binary analyses between each group pair (see [Aly, 2005](#) for an overview of the main types introduced below).

Selecting this option will reveal the following menu:

- 1 : Simple Pairwise Decoding (Maximum-Wins method)
- 2 : Error-correcting output codes
- 3 : Hierarchical One-vs-One method

Simple pairwise decoding This is the simplest method. Given that NeuroMiner has conducted different pair-wise classifications between all the group pairs, this option will categorise the individual by their maximum score across all the classifiers. For example, in a three group problem, if an individual's score is highest for the schizophrenia group in the comparison between schizophrenia and bipolar (i.e., as opposed to the SCZ-BP, or SCZ-DEP, or BP-DEP) then they will be classified in the schizophrenia group.

Error-correcting output codes Error-correcting output codes (ECOC) is a meta-method that is used to solve multi-group classification problems that applies a voting scheme (i.e., an ensemble method) to decide upon the correct class. It works by assigning a binary code to each group and then binary functions are learned one for each bit position in the binary code strings (see [Dietterich et al., 1995](#) and [Dietterich and Bakiri, 1991](#) and [Kinderman et al.](#) and for an overview see [Aly, 2005](#)). Optimal results are obtained by maximally separating the codes using a distance metric, such as the Hamming distance.

Hierarchical One-vs-One Method This technique arranges the classes into a tree and then at each node of the tree a binary classifier makes the discrimination between the different child class clusters (see [Aly, 2005](#) and [Wang et al., 2009](#)).

4.4 Stacking

Stacking is an advanced feature of NeuroMiner that takes multiple learned decision values, combines them, and then learns again from this meta-information (see [wiki](#)). This technique has provided better solutions than using single models to learn the classification or regression problem. Stacking can be done in NeuroMiner if you have different data modalities (e.g., you can learn from clinical and neuroimaging separately and then stack the learners) or if you have run separate analyses on one (or more) modalities (e.g., if you run a linear SVM and then a non-linear SVM and then combine the learned models).

If you have more than one modality or have run separate analyses on the same data, then an option will appear in the "Define parameter template" menu entitled "1: Define meta-learning/stacking options".

- 1 : Activating stacking using prediction generated by other analyses []
- 2 : Select 1st-level analyses to provide features to stacker [selected:]
- 3 : Define feature names [] 4 : Extract predictions from ...

1: Activating stacking using prediction generated by other analyses Turn stacking on or off.

2 : Select 1st-level analyses to provide features to stacker Of the analyses that have been conducted, select the models that you want to stack.

3 : Define feature names [Var 'feats'] Optionally define feature names for each of the models that you want to stack.

4 : Extract predictions from ... This gives the option to extract predictions from either the CV1 training data or the CV1 test data (the default). If CV1 test data is used, the higher-level algorithms should not be combined with wrappers because of overfitting. Cross-CV probabilistic feature extraction is helpful in this regard to prevent overfitting.

Once stacking is turned on and these options are defined, then all parameter settings now relate to the stacked models (e.g., all the preprocessing and training settings). These settings will need to be changed because the data inputs are now the decision scores for each of the models (i.e., you'll have to change the preprocessing accordingly). Then you can return to the main menu, initialize the stacking models, and train as normal. The results will then relate to the stacked models rather than the original data.

It's important to note that the results now pertain to the decision scores of the stacked models and therefore the interpretation of results is based on the

models and not the original data (i.e., you can see how one model performed over another, but not information related to the individual variables).

Please note, that stacking does not work with external validation data.

4.5 Initialize analyses

→ After defining analyses you must initialise them

As outlined in section 4, after the machine learning parameters have been defined in the parameter template (section 4.2) then you have to initialize them for further analysis using the initialisation manager. If analyses in the parameter template are changed, you must also re-initialise them using this menu. You can initialise and change as many analyses as you like.

When an analysis is first initialised, a directory will be created containing details and parameters of the analysis in a log file (see below). During processing, NeuroMiner will store analysis files in this location for this analysis. These data reporting features allow you to keep track of analyses and provide a standardised NeuroMiner output format.

It's important to note that within NeuroMiner you can initialise one analysis, then change settings in the parameter template (see 4.2), and then initialize another analysis. This has two implications: 1) the results from different analyses can be compared; 2) with the results of different classifiers, you can use the advanced feature called stacked generalization (see 4.4). It is important to note that simply changing settings in the parameter template does not have any effect on analyses until you initialise the analyses using this menu. If you want to change the settings of an analysis then change them in the parameter template and then reset the initialised analyses using the directions below.

The initialization menu can be split into two parts: 1) when no other analysis is present; 2) when there are existing analyses. And will result in the outputs described below.

4.5.1 First Initialization

The following menu will appear:

```
Define analysis identifier [ ? ]  
Provide analysis description [ ? ]  
Specify parent directory of the analysis [ ? ]  
PROCEED >>>
```

Define analysis identifier Provide a brief description of the analysis; e.g., SVM_analysis1. This will be used as an analysis ID and also used as an identifier for the analysis-specific directory created by NeuroMiner.

Provide analysis description Here you can write a more extended analysis description; e.g., "this is a test analysis that was created using the default settings of NeuroMiner". You can enter multiple lines if you would like and then enter a single period (i.e., ".") to stop the process and enter the data.

Specify parent directory of the analysis This will activate a menu selector where you can specify a parent directory for the analysis folder to be stored.

PROCEED >>> This will activate the above settings.

4.5.2 Multiple Analyses

If an analysis is already initialized, then entering the initialization menu will show the following:

```
Generate new or manage existing analyses [ generate new: analysis 2 ]
Define analysis identifier [ ? ]
Provide analysis description [ ? ]
Specify parent directory of the analysis [ ? ]
PROCEED >>>
```

NeuroMiner defaults to generating a new analysis (option 1), and new data for the aforementioned fields can be entered that were described above. Once this is completed, a new analysis directory will be established and you can proceed as normal. However, when there are multiple analyses, there is also the possibility of managing the existing analyses by selecting the first option, then selecting "manage existing", which will show the following menu:

```
Generate new or manage existing analyses [ manage existing ]
Select existing analysis [ Analysis 1 ]
Specify what to do with selected analysis [ complete reset ]
Define analysis identifier [ test1 ]
Provide analysis description [ provided ]
Specify parent directory of the analysis [ ]
PROCEED >>>
```

The second option will allow you to select the analysis. If there are two or more analyses, then it will activate a "Select Analysis Menu" with the parameters at the top and then an analysis manager at the bottom. For example:

MODALITY 1: MATLAB matrix

Preprocessing:

Group processing mode: binary

Step 1: Scale [from 0 to 1], zero-out completely non-finite features

Step 2: Prune non-informative columns from matrix [Zero Var, Inf]

Step 3: Imputation in Matrix Block [All features]: median of 7 nearest neighbors (Euclidean)

Machine Learning Method:

GLM, Linear Regression Model, No Kernel

No optimization of preprocessing parameters

No optimization of ML parameters

```
Cross-Validation:  
CV2: 10x10, CV1: 10x10  
  
ANALYSIS MANAGER: CURRENT ANALYSIS: 2 (of 2 ANALYSES)  
  
1 : Select current analysis  
2 : Generate new analysis  
3 : <= Go to previous analysis  
<== Back/Quit [Q]
```

This menu will give the user the ability to view what has been initialised and to navigate through the analyses and choose the right one.

Once an analysis has been selected, the user will be returned to the main menu and can then modify the analysis by selecting the third option to "Specify what to do with the analysis":

- 1 : delete
- 2 : completely reset
- 3 : reset parameters (risk of inconsistency)

Here the user can select the option to delete the analysis, to completely reset and overwrite all the parameter settings with the new settings in the parameter template, or to reset parameters. Once an option is selected, you'll be taken back to the main analysis manager and you will need to "PROCEED >>>" for the options to take effect.

4.5.3 Outputs and Directory Structure

Once an analysis has been initialised, a folder will be created like the following:

NM_ID736740_54131_A1_your-description

This outlines the NM identifier associated with the entered data (e.g., ID736740_54131), the analysis that has been initialised (e.g., _A1_), and the description that you entered in the "Define analysis identifier".

Within the folder there is a "params" subfolder and also a log file (e.g., "NM_Analysis_test1.log"). The log file contains an "NM ANALYSIS LOG FILE" with a description of the computer and the versions of MATLAB and NeuroMiner, in addition to a description of all the analysis parameters that were conducted. The params subfolder has these details and more specific settings stored in a .mat file containing MATLAB structures that are used by NeuroMiner.

Following initialization, you can then go to the main menu and preprocess or train the

4.5.4 Once an Analysis has been Initialized

Once an analysis has been initialised, you can then preprocess the features and train the models using the directions below. These files will be stored in

the directory for that analysis.

4.6 Preprocess features

Once the data has been initialized, the user has the option to preprocess the features. In NeuroMiner nomenclature, this refers to any steps that are conducted within the cross-validation folds prior to training a model. For example, the data within each fold could be scaled and then imputed, or a principal components analysis could be conducted.

When this option is selected, the user will be presented with the following:

- 1 : Select analysis to operate on [Analysis 1]
- 2 : Overwrite existing preprocessed data files [no]
- 3 : Use existing preprocessing parameter files, if available [no]
- 4 : Write out computed parameters to disk (may require A LOT of disk space)
[no]
- 5 : Select CV2 partitions to operate on [0 selected]

2: Overwrite existing preprocessed data files. NeuroMiner will save pre-processed data files and selecting this option will overwrite them.

3: Use existing preprocessing parameter files, if available. The user has the option to save parameter files (e.g., betas from regression of age from a feature) and this option will allow them to use the existing files.

4: Write out computer parameters to disk. This option will write the parameters (e.g., beta coefficients) to disk. This is not recommended due to space limitations.

5: Select CV2 partition to operate on. This option will only appear if an analysis has been selected in step 1. It allows the user to select the folds that the preprocessing analysis will be applied to (see Box: Selecting CV folds in NeuroMiner).

Box: Selecting CV folds in NeuroMiner. Because NeuroMiner was developed with nested, double cross-validation as the gold standard, there is an option to only select the outer, CV2, folds that you want to process. For example, you could select only the first permutation in order to check that everything is working before running the rest of the permutations.

When the option is selected, the user will see the CV2 cross-validation folds in a block of permutations and folds:

e.g., Perm 1: *-*-*-*-. A star "*" indicates that the fold **is selected**, whereas the letter "o" indicates the fold **is not** selected; e.g., Perm 1: o_o_o_o_o. Please note that all folds are selected by default in NeuroMiner, and thus you would need to de-select all folds first if you wanted to select specific perms or folds. The user can select and de-select the folds in single positions or ranges. The folds can be referenced by first indicating the permutation (e.g., 1st permutation) and then the fold position (e.g., 1st fold). For example, the second fold of the first permutation would be [1,2]. A fold range can also be selected in the following format [perm start, perm end, fold start, fold end]. For example, you may want to select the first two permutations and the first three folds in the following way: [1,2,1,3].

4.7 Train supervised classifiers

As depicted in Figure 3, once the parameter template has been completed and initialized there is the option to train the models. This implies the use of a statistical technique, such as an SVM, along with the hyperparameter optimisation (e.g., optimising the C parameter) and any wrappers that might have been activated. The following menu will be displayed:

- 1: Create analysis from scratch
- 2: Create analysis using precomputer preprocdata-MATs
- 3: Create analysis using precomputed CVdatamats
- 4: Create analysis using precomputed CVresult-MATs

1: Create analysis from scratch

This is where the analysis has only been initialized and no data matrices have been stored. This will preprocess and then train classifiers. Please note, that if your matrices are large (e.g., with neuroimaging data) then this will take up a lot of RAM and as such it is recommended to preprocess the data first.

2: Create analysis using precomputer preprocdata-MATs

This is when the user has preprocessed their data and stored the preprocessed files in a folder; i.e., they have TRAIN files.

3: Create analysis using precomputed CVdatamats

This is when the user has completed an analysis (preprocessing and training) and they have the data matrix files stored in a folder; i.e., "...CVdatamat_oCV..." files. When this option is selected, the analysis results from the CVdatamat files are then collated into a result file. It is useful especially when the program

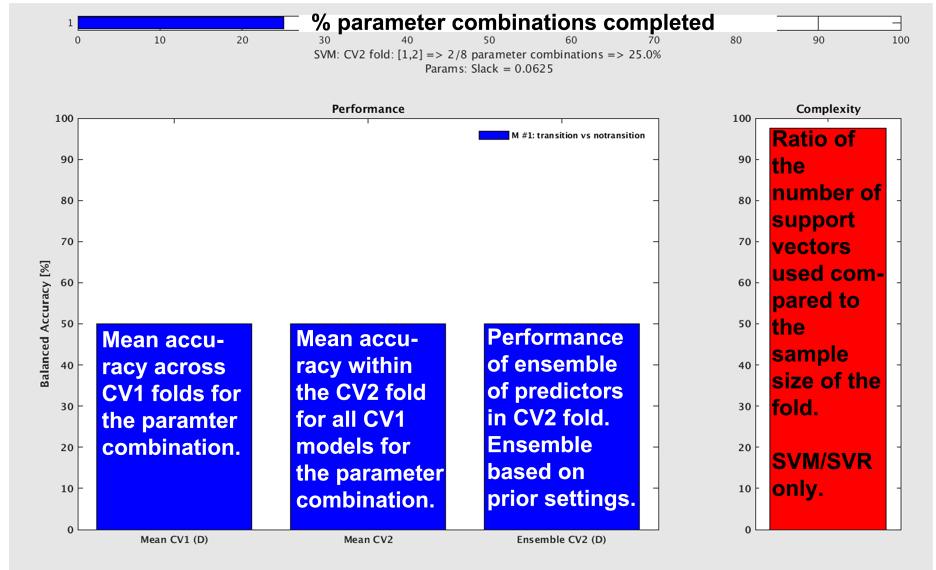


Figure 6: Monitor progress during an analysis.

has crashed, and/or the NM structure has not been saved, and the user wants to get the results without re-running all the training again.

4: Create analysis using precomputed CVresult-MATs

This is when the user already has the results file and they want to integrate the results into the NM structure. This is useful if there has been a crash and the data has been lost from the NM structure.

4.8 Visualize classifiers

A standard NeuroMiner analysis will produce the predictive accuracy of the models using the preprocessed data (e.g., the scaled PCA components). This means that the model weights (e.g., from an SVM) are only available within this space (e.g., a weight for each PCA component). For accurate interpretation, it is necessary to project these weights back to the original data and this is called "visualisation" within NeuroMiner. The term visualisation applies best to NeuroImaging data where the weights are projected back to the original voxel space and then can be visualized, but also applies to matrix data where the results for each test (e.g., for each questionnaire item) can be visualized in a graph. Running the visualization routine in NeuroMiner will produce three different types of results that represent how the classifier relates to the original data: mean results, grand mean results, and z/p value results associated with the individual feature weights.

Visualisation in NeuroMiner also applies to the calculation of the model-based permutation p-values that can be selected in the "Visualisation" menu.

Mean Results

****_Mean_****

This image represents the median of the weights across all models selected during the CV1 cycle. For example, if you have a 10x10 cross-validation cycle and you're selecting the winner for each fold, then it will create a median across 100 models.

****_SE_****

The standard error of the weights across all models selected during the CV1 cycle.

****_CVratio_****

This is the sum of the median weights across all CV1 folds divided by the standard error. It is a measure of stability.

Grand Mean Results

****_Mean-GrM_****

Within a nested cross-validation framework, the grand median (GrM) in NeuroMiner is defined as the sum of the median weights of selected CV1 models for each CV2 fold, divided by the number of CV2 folds.

****_SE-GrM_****

The grand median standard error is the sum of the standard error values of selected CV1 models for each CV2 fold, divided by the number of CV2 folds.

****_CVratio-GrM_****

The grand median CV ratio is the sum across CV2 folds of the selected CV1 median weights divided by the selected CV1 standard error. This value is then divided by the number of CV2 folds.

****_Mean-gr-SE-GrM_****

Across selected models for CV2 folds, this is the grand median thresholded to only display the voxels where the median is greater than the standard error.

****_Prob95CI-GrM_****

Across selected models for CV2 folds, this is the amount of times that the median value has been greater or less than the 95% confidence interval.

****_Sign-Based-Consistency_****

Sign-based consistency mapping for features is based on the logic outlined in [Gomez-Verdejo et al., 2019](#). The importance is calculated as the number of times that the sign of the feature (i.e., a positive "+" or negative "-" sign) is consistent within an ensemble (e.g., number of times that it is consistently positive) multiplied by the number of times that the variable was non-zero (e.g.,

in the case of a L1-regularized SVM or regression, the weights can be set to zero and thus the more that this happens the lower the importance of the feature). The measure is between 0 to 1, with 1 representing perfect consistency within the ensemble and 0 if the weights are equally positive and negative or when the feature is omitted with a zero weight. A p-value is then calculated by defining a hypothesis test for the importance score with a null hypothesis of 0 and an alternative hypothesis of $\neq 0$. A z-score was calculated as the importance divided by the square root of the variance of the importance scores. A standard p-value was then calculated using a normal cumulative distribution function to choose the right-tailed significance. P-values are then corrected using the false-discovery rate.

Z and P score images Visualisation can also include the derivation of Z or P scores for the entire models (i.e., not the individual features) using permutation testing. These are outlined in [4.2.7](#).

4.9 Display training results

Once initialization, preprocessing, and model training has been performed, the results can then be displayed using this option. The results display interface can be understood in three main divisions: 1) the accuracy and model performance (initial screen); 2) the model performance across the selected hyperparameters; 3) results concerning the individual features in the analysis.

4.9.0.1 Accuracy and Model Performance See Fig. [7](#). The performance of a model can be fundamentally described by the degree to which the predictions match the target scores (e.g., how many people are classified correctly or how closely the features and target are related in a regression). In NeuroMiner, this is represented in multiple ways that each address a different aspect of the relationship between subjects, predictions, and targets.

4.9.0.2 Model Comparisons See Fig.[14](#). When there is more than one analysis, statistical comparisons can be conducted between the model performances using a variety of statistics using the "Statistical Comparisons" option. The results of these comparisons are output in either excel spreadsheet files (i.e., xlsx) or in text files (i.e., csv) for Linux and Mac. The menu is quite intuitive (i.e., as below) where you select the analyses you want to compare on the right, you press the "Save As" button to choose a name, and then you select the "Compare" button to generate the files. The files will be generated in your working directory.

See Fig. [15](#). The user can also perform "Visual Comparisons". The user needs to select some performance measures that they want to compare visually in the top-half of the pop-up box using the "Select" column tick-boxes, then selects the analyses in the lower-half of the pop-up box. Then simply press "Visualise".

If the bar plots are not good then the user can change it to "Line plots" and then press "Visualise" again.

4.9.0.3 Classification Performance See Fig. 8. NeuroMiner was built to work with hyperparameter optimisation within a nested-cross validation framework as a standard. In this plot, you will see the inner-cycle, CV1, mean test performance plotted in blue and then the outer cycle, CV2, mean test performance plotted in red. Each point represents a hyperparameter combination. This gives an indication of how the hyperparameter combinations are influencing the test performance. On the right, the display also shows heatmaps that show how the performance in the CV1/CV2 folds changes over the permutations.

4.9.0.4 Generalization Error See Fig. 9. Another way to look at CV2 test performance is with the generalization error, which is calculated in NeuroMiner as the CV1-test performance minus the CV2 test performance.

4.9.0.5 Model Complexity See Fig. 10. In NeuroMiner, complexity is defined as the percentage of support vectors divided by the number of subjects in the tested fold. This feature is useful to determine how varying parameter combinations may affect the fit of the model.

4.9.0.6 Ensemble Diversity See Fig. 11. This is a measure of Shannon entropy across CV1 models for each CV2 partition. A higher entropy means higher disagreement between the CV1 models. This measure is useful because increased entropy will often result in a better fitting model at the CV2 level. It is useful to note that entropy can be maximised as a performance criterion in addition to accuracy.

4.9.0.7 Model Selection Frequency See Fig. 12. Within an ensemble framework, the models that are selected within the inner, CV1, cycle are then applied to the outer test data. For example, if it is a 10x10 CV1 framework with a winner-takes-all method of model selection, then 100 models will be applied to the held-out CV2 test fold. This plot displays the percentage of times that each parameter combination was selected.

4.9.0.8 Multi-Group Results If a multi-group analysis is conducted then the results display will change to accommodate the groups. See figure 13.

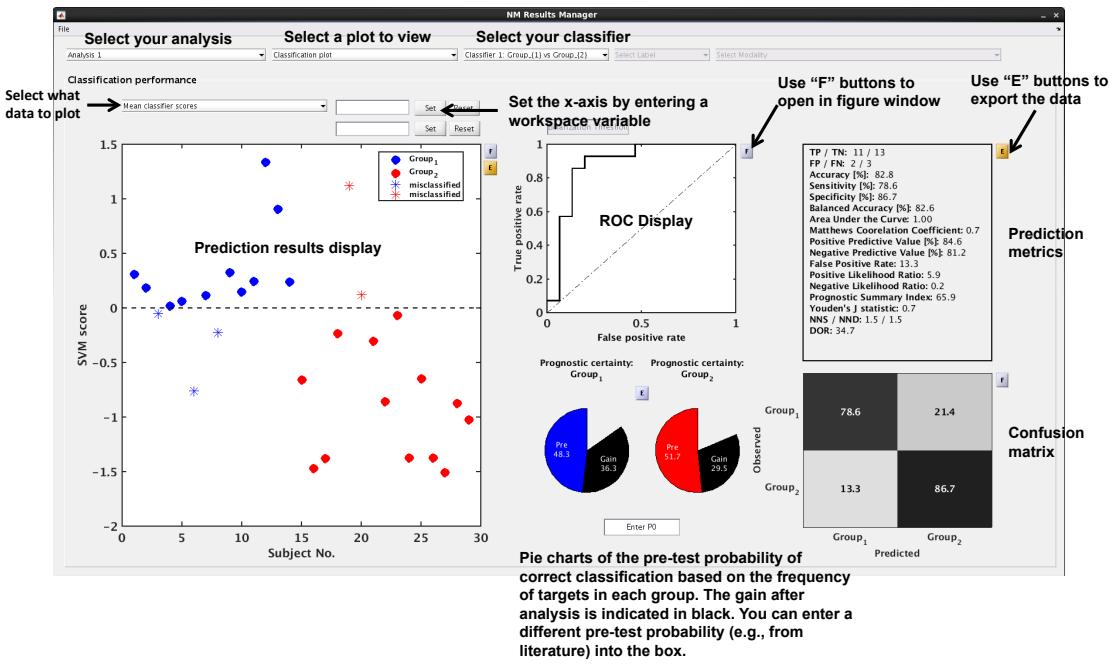


Figure 7: Option 1: classification plot

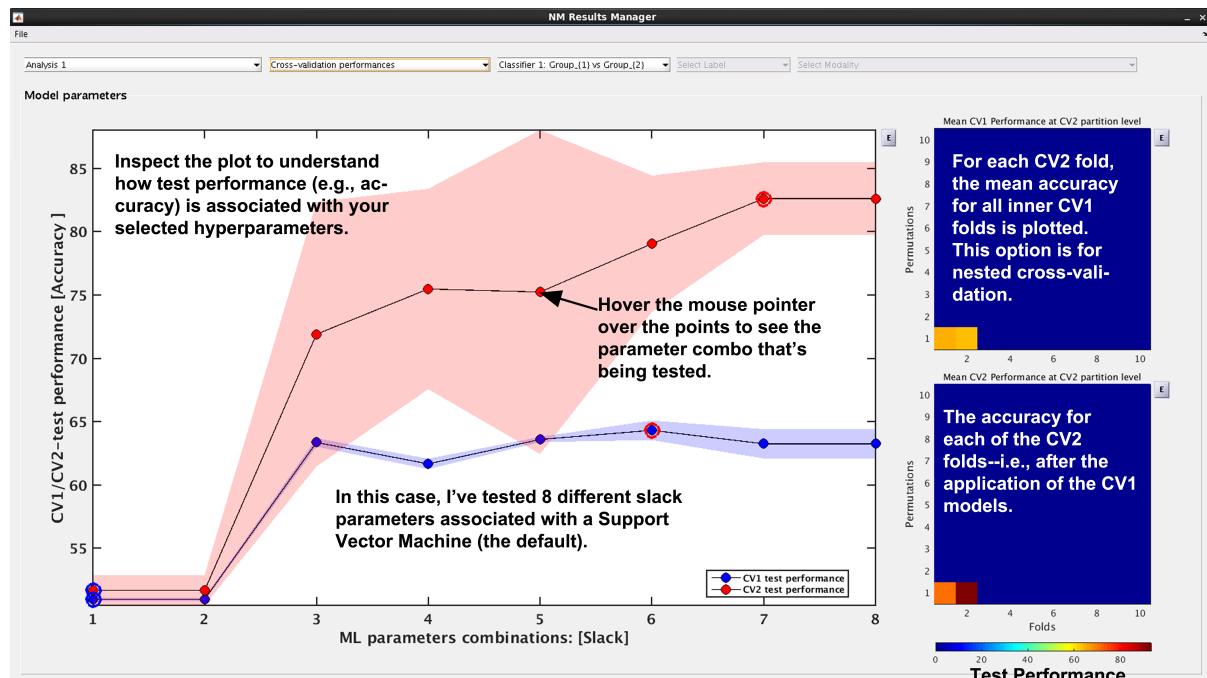


Figure 8: Option 2: Cross-validation Performance.

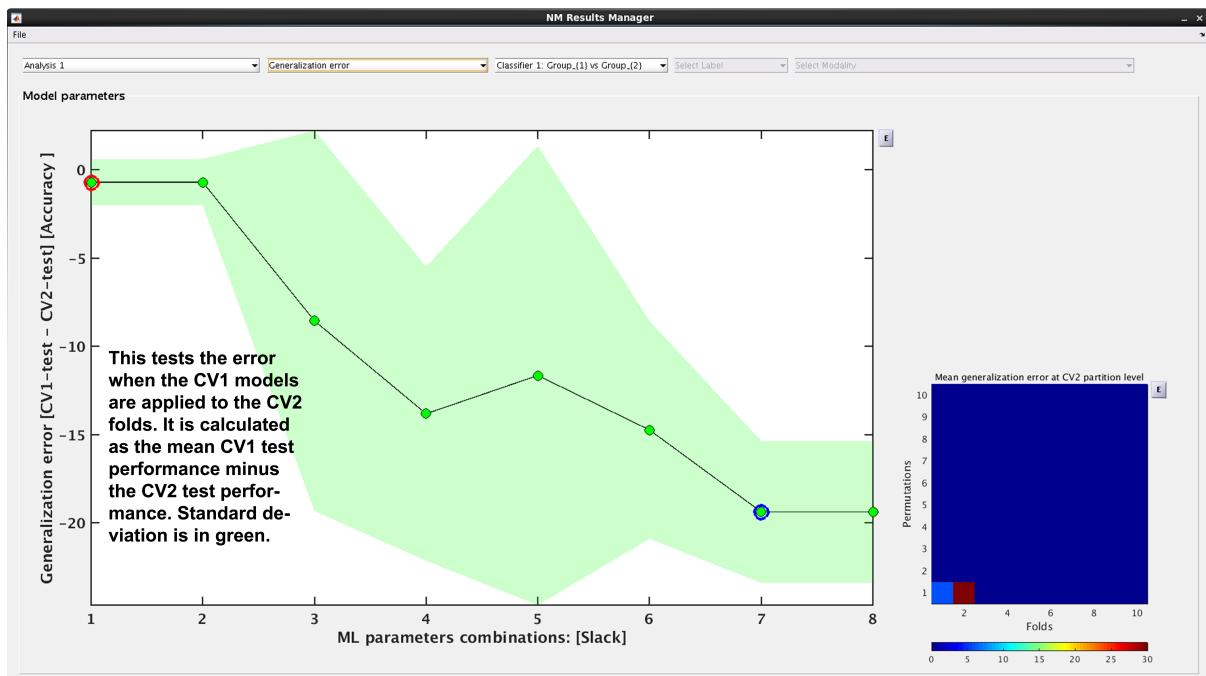


Figure 9: Option 3: Generalization Error.

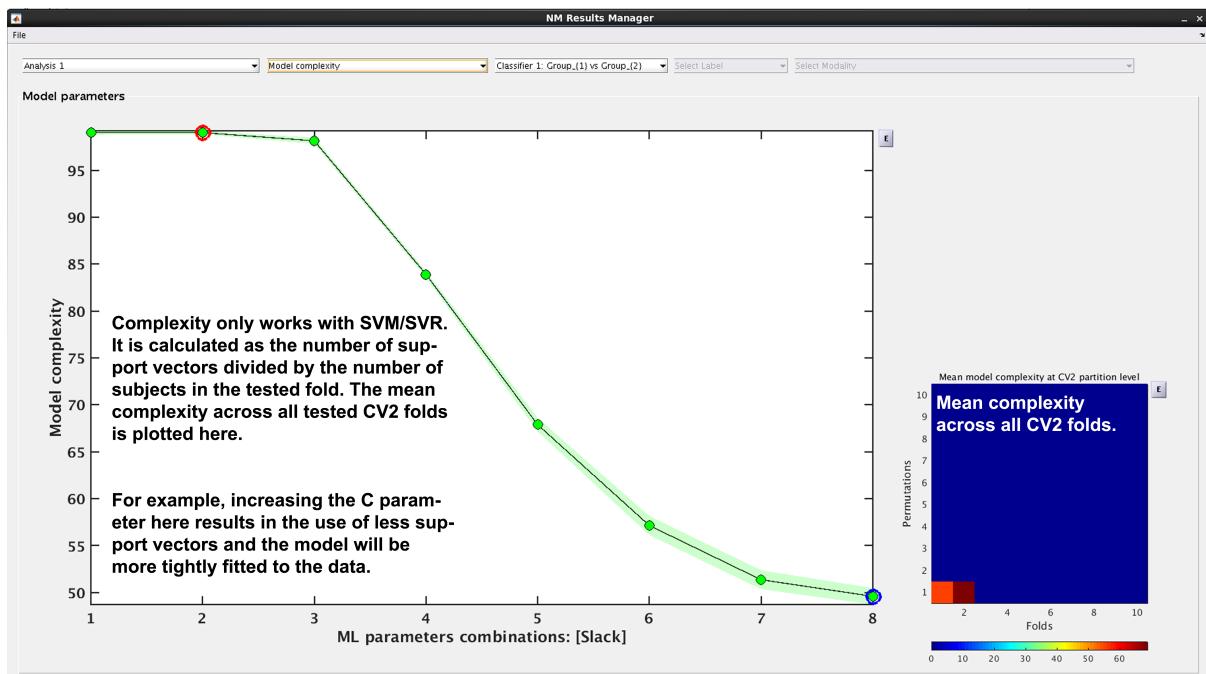


Figure 10: Option 4: Model Complexity

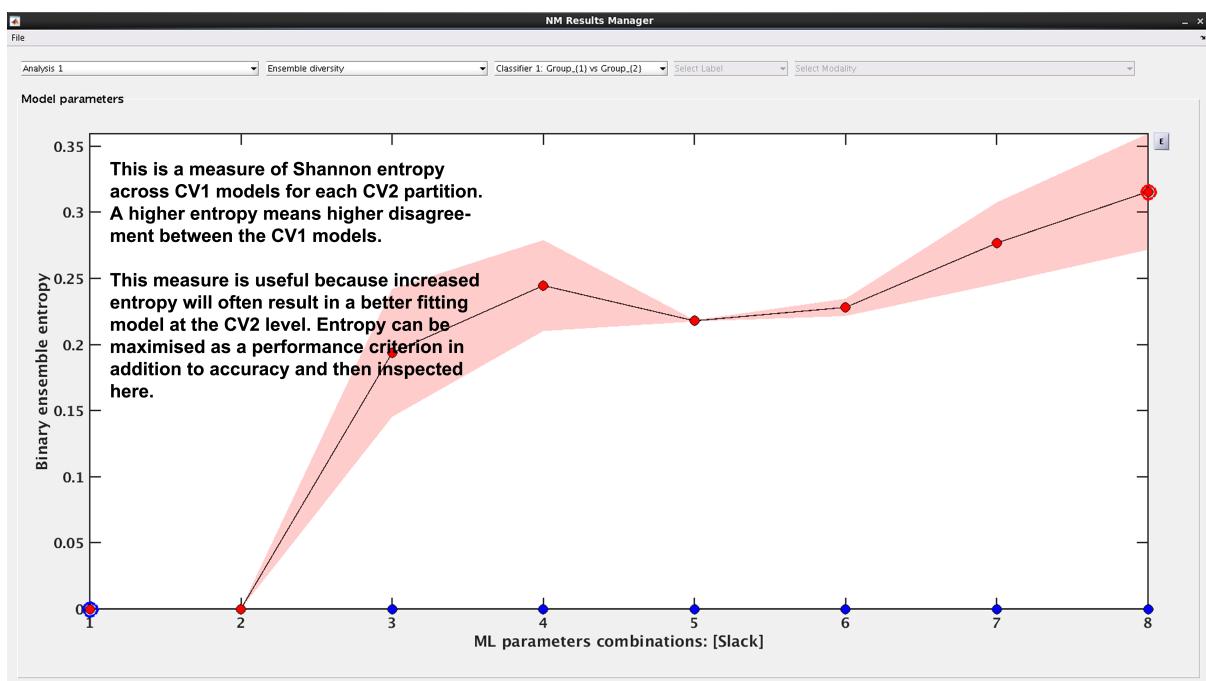


Figure 11: Option 4: Ensemble Diversity

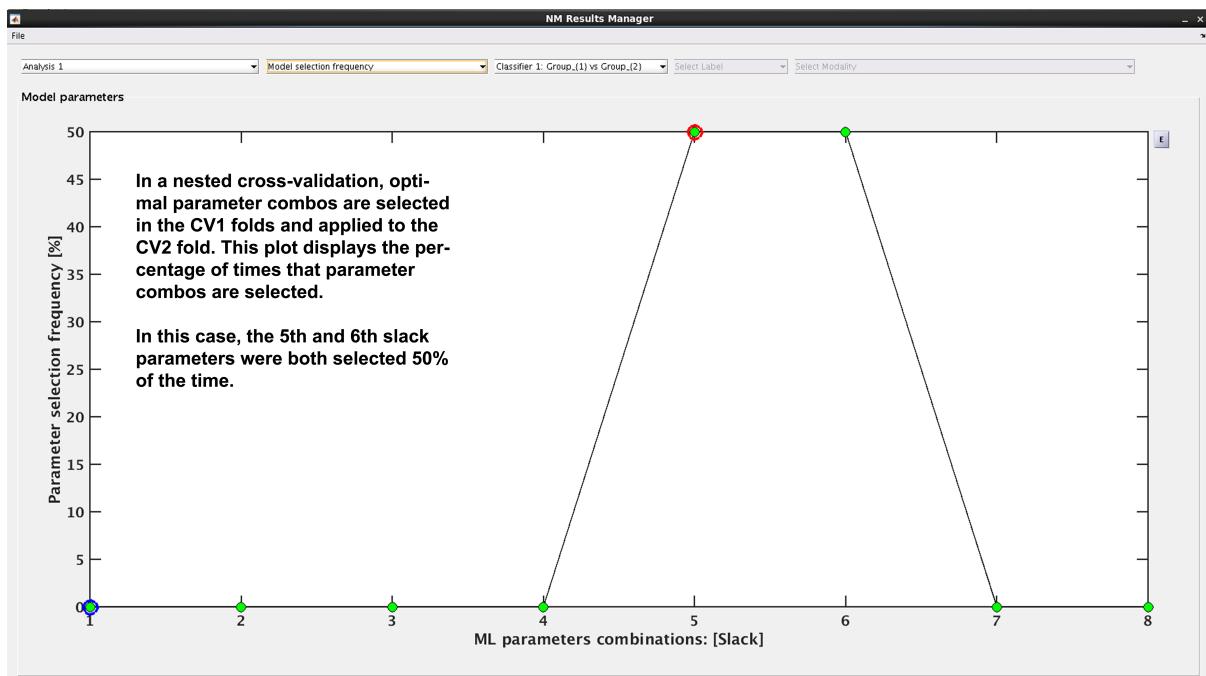


Figure 12: Option 5: Model Selection Frequency

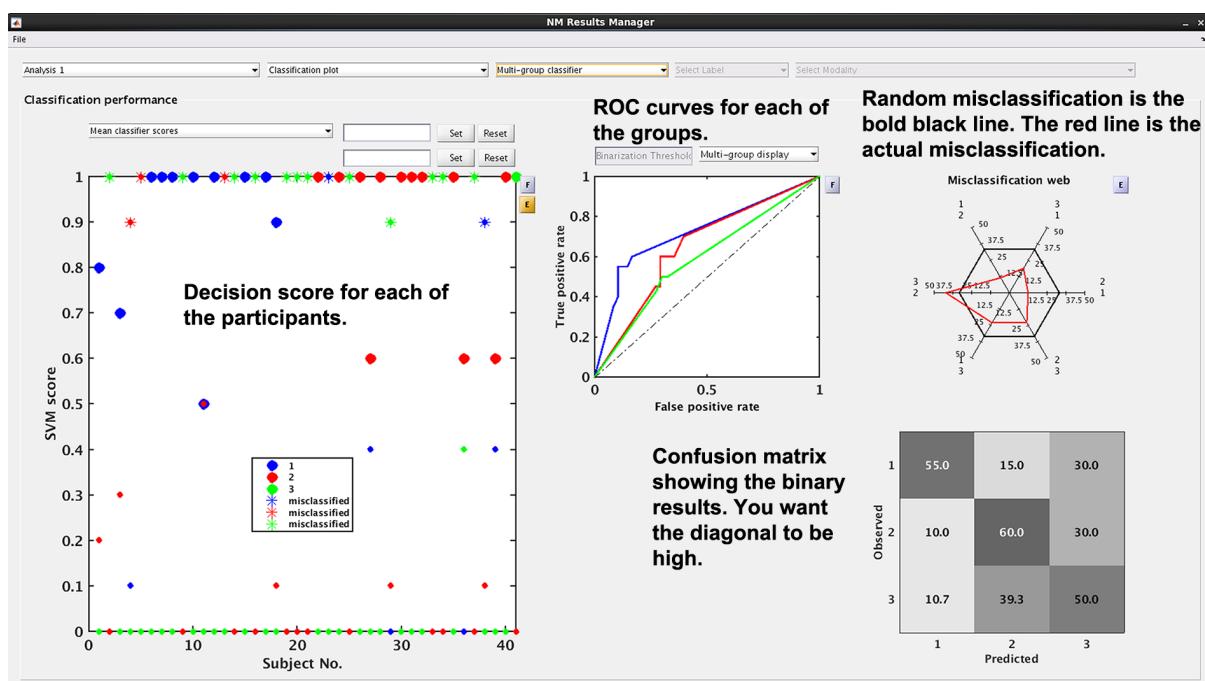


Figure 13: The multi-group display

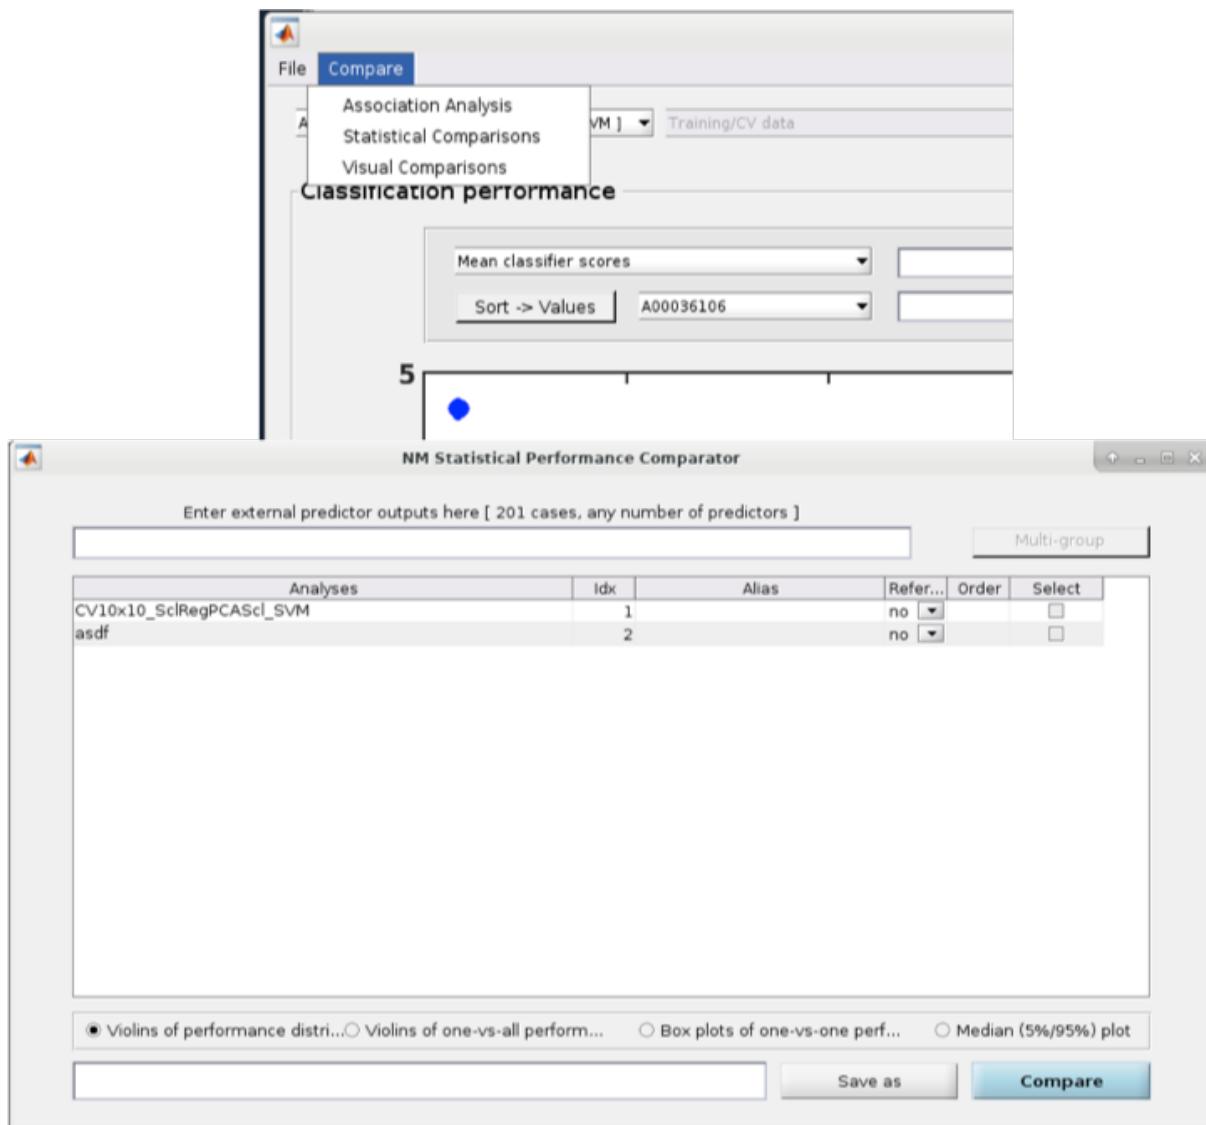


Figure 14: Within the main results display, the user can go to the "Comparisons" menu item and then select the option to calculate statistical differences between separate models

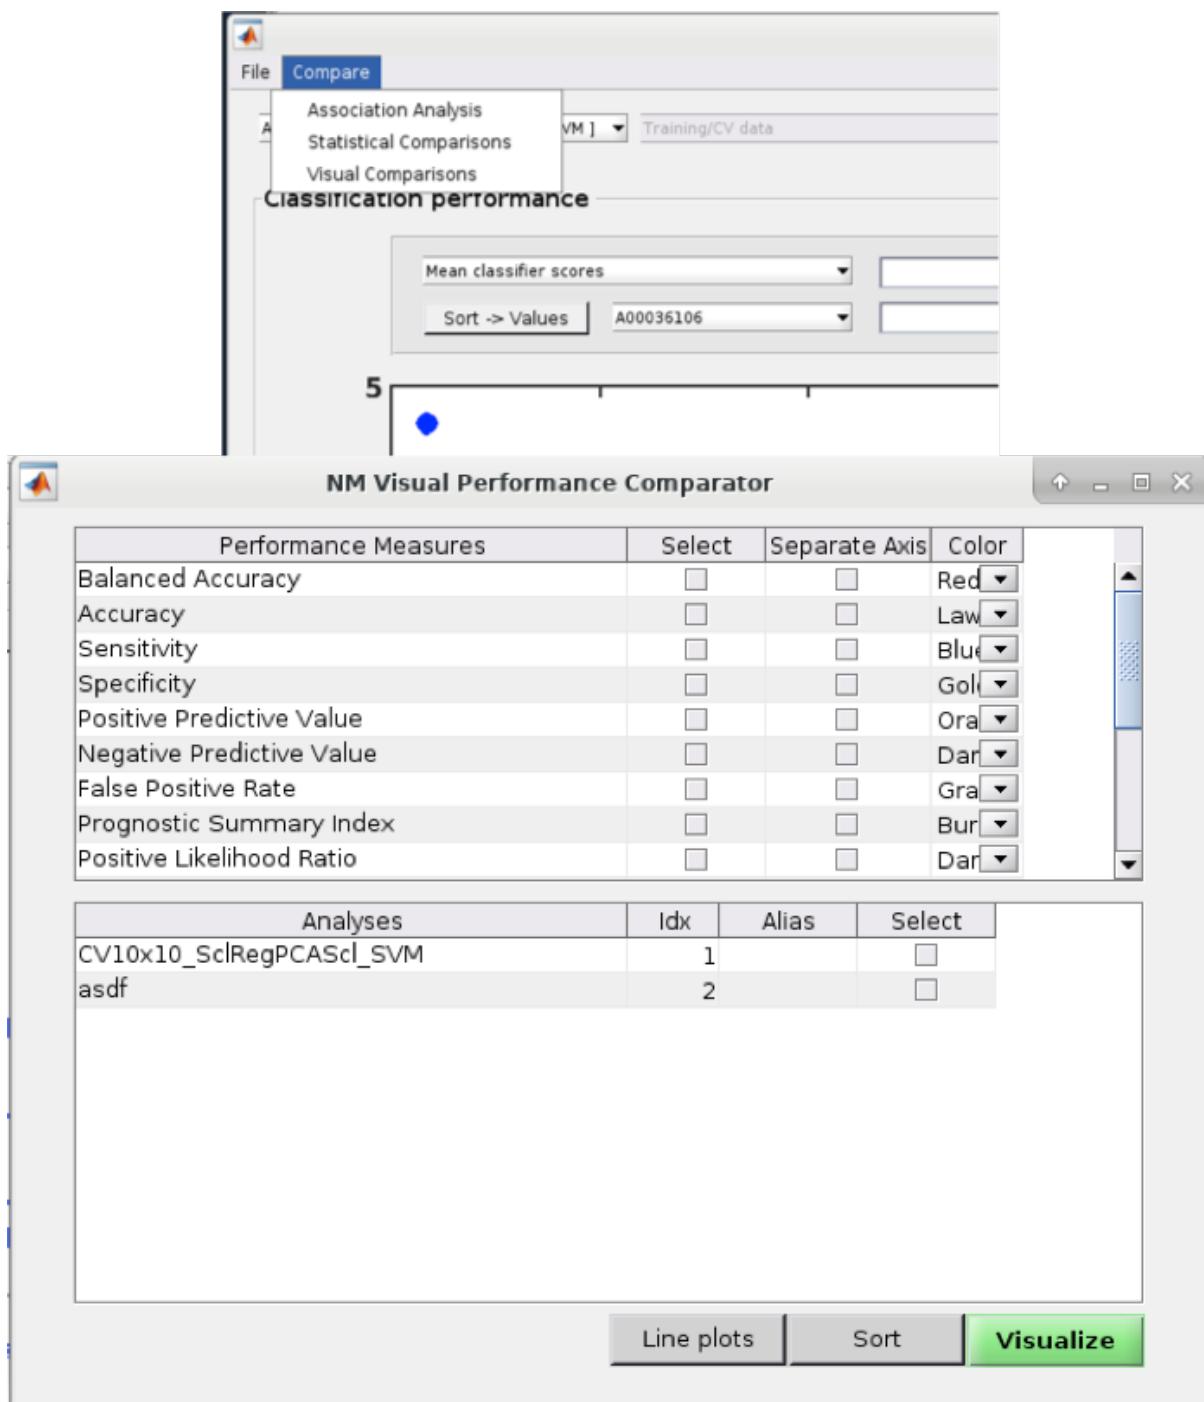


Figure 15: Within the main results display, the user can go to the "Comparisons" menu item and then select the option to display visual comparisons between separate models

4.10 Generate master files

Master files of the individual preprocessing or analysis files can be created using this option.

4.11 Out of Sample Cross-Validation

External, out-of-sample cross-validation (OOCV) is a cornerstone of machine learning and is required for strong generalizability claims. This usually involves the application of the trained models used to predict targets in the CV2 data folds to a separate dataset that has been independently collected.

External validation relies on the trained models being completely separate to the OOCV data. For this reason, after completing the analysis the user must first lock analyses and start the NM application mode using the following procedure. As mentioned in the section on cross-validation, a form of external validation can be done within a machine learning analysis by selecting the leave-group-out option if the group comes from an external source (i.e., it is not from the same study or sample). However, mostly external validation is done when models from one analysis are applied to a completely different site and study.

It is important to note that within a machine learning framework external validation never means that new models are created within a different sample and then the results are compared (e.g., comparison of the beta coefficients or the p-value for significance). This is usually referred to as replication and is basically never the point of machine learning because it is a field that grew around translation and practical application of models. The aim of machine learning is to create models that can be applied without modification to external samples to facilitate predictions at another site. The ultimate aim is to be able to integrate models into clinical care globally.

For this critical function, what needs to occur is that the models that were applied to the test sets at the outer CV2 level need to be applied to another sample. Keep in-mind that when we refer to models in NeuroMiner then we are referring to the entire pipeline of operations that are applied to the data: from preprocessing to training. This can be done in multiple ways, depending on the method of calculating each model in the training data. In NeuroMiner, as described in the section on cross-validation, the optimal models are ultimately retrained in each of the CV1 folds before being applied to the outer test folds. This means that the same models will be applied to the external validation data. For example, if there are 5 folds on the outer cycle and 5 folds on the inner cycle then $5 \times 5 = 25$ model pipelines will be applied to the external validation data (assuming that the single best model is chosen for each training fold).

As stated in the section on cross-validation, in order to get a valid measure of generalizability it is critical that all operations included in the model training process are applied without modification to the external validation. We can also

add that it is also critical that different model sets aren't repeatedly applied to the external validation data so that the user can choose what parameters and options fit the external data the best this would completely undermine the reason for performing external validation. For this reason, in NeuroMiner features have been developed so that the user can't easily re-run analyses with the external validation.

We will be using external validation functions for the machine learning competition in order to test the accuracy of the models you develop to a completely different testing set that has come from a different study and scanner. To do this, we need to load a NeuroMiner structure where there is an analysis or analyses that you would like to apply to the external validation sample and get the external validation data.

Once an analysis has been completed, from the MAIN INTERFACE, choose:

8: Lock analyses and start NM application mode.

Then select "y" to close the model discovery.

The MAIN INTERFACE will then change to include:

1: Load data for model application.

Select this option.

A dialogue box will appear (Fig 16). To enter new data, the user first needs to know whether the Target labels are known (i.e., do you know what group the cases belong to if it is a classification?). If they are known then select the tick-box at the bottom right called Target labels known and if they are not (e.g., for the final competition submission) then do not select this box. The same procedure applies that is described in 4.1. Then select the Create New box at the bottom right and it will create a field in the middle blue center box that looks like this:

1: (0 cases); date: DATE; labels known: no

This is an empty template that needs to be filled with data by pressing the box at the top right called Add/Modify Data in Dataset and following this you will be taken to the data entry menu of NeuroMiner, which contains some of the same input settings as described in the original data entry exercise we completed at the start of the school. It is important to note that you cannot move the original data from its original location (including any masks used in the processing) otherwise NeuroMiner will fail at this step. You should be able to see the following interface like the following for 'Input data into NM':

1: Select images files [?]

2: IMPORT nifti

A point to note that the type of Input data has to be same as that of the original data type used. NeuroMiner automatically detects the previous input data type

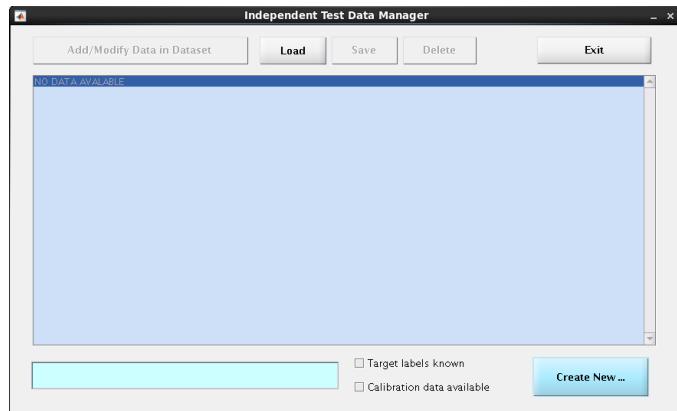


Figure 16: The independent test data dialogue box. 1) Press "Create New...";
2) Click on the new dataset line; 3) Press "Add/Modify Data in Dataset".

and asks accordingly. For example if the initially entered data was a 'Spreadsheet like data', the data to be entered in the model application mode has to be the same data type. If the user has entered covariates during the discovery analysis then NeuroMiner will ask for a covariates variable. Covariates always need to be entered into the matlab workspace as a variable (i.e., similar to when we were doing data entry for the discovery sample). Once, everything is setup, you should press:

6: IMPORT nifti

Once the data has been entered, you can return to the MAIN INTERFACE and now there will be an option

3: Set up parameters for model application

Select this option to see the options for defining the parameters in OOCV settings:

1: Aggregate all base learners into one big ensemble

2: Model retraining mode [as defined for CV model training]

The first main option is 'Ensemble mode' which relates to an important way that the models are applied to the external validation sample that is relevant to all machine learning analyses. Upon entering this option, you would see the following two options:

1: Aggregate ALL base learners into ONE ensemble

2: Compute mean decisions of CV1 partition ensembles before aggregating

When you have a cross-validation design, it means that all models relate to

one of the outer, CV2, folds. However, when you apply the models to a new sample, generally all the models will be applied to this sample (this is the Aggregate ALL base learners into ONE ensemble option below). This means that the procedure applied to the external validation sample is different to what is applied to the discovery sample. For this reason, NeuroMiner provides the option to apply each CV1 partition ensemble for each CV2 fold before computing a mean of these decisions in order to follow a more similar logic to the discovery sample procedure.

The next important option is the 2: Model retraining mode option. The two sub-options on entering it are:

- 1:As defined for CV training phase**
- 2: Use all available CV data (CV1+CV2)**

This is a critical aspect of NeuroMiner that needs to be understood from a machine learning perspective. As a default, NeuroMiner does not save the specific model parameters to disk because they generally consume a large amount of disk space (e.g., 50G). Instead, NeuroMiner defines the optimal parameters for the models during the preprocessing and training procedures. It then retrains the models in the discovery sample using these parameters before applying them to the external validation sample. For this option, NeuroMiner gives the possibility to retrain the models in the same way that they were trained on the discovery sample in 1 or to 'Use all available CV data (CV1+CV2)'. The first option is recommended if the point of the analysis is to validate the original models that will be displayed in a paper. The gold standard though would be to save the models and use the exact models calculated in the discovery cohort for the external validation cohort, which can be done during Visualisation in NeuroMiner. The reason for this is that there can be minor changes to the models when they use non-deterministic procedures (e.g., gradient descent).

After setting up the parameters, simply exiting to the MAIN INTERFACE in order to accept the defaults here. In the MAIN INTERFACE, you will now see an option of

4: Apply classifiers to independent data

The user should select this option. Using the winning parameters from your discovery analyses, the models will now be retrained in the discovery data before being applied to the external validation data. You can see this happening in the matlab command window. After it has been completed, you can then go back to the MAIN INTERFACE and select:

5: Open NM Results Viewer

You will then be taken to the Results Display as before. If you click on the sec-

ond drop-down box then there will be an option to select the external validation data that you entered.

4.12 Load NeuroMiner structure

As outlined in section 4, the NM structure is the core structure that NeuroMiner uses to store and process data. This option allows the user to load it from a .mat file.

4.13 Save NeuroMiner structure

This option allows the user to save the NeuroMiner structure (NM) as a .mat file in a folder of the users choosing. Notably, any changes that are made to the settings of NeuroMiner during an active session are stored within the NM structure in the workspace. However, program crashes can lead to the loss of information. We recommend to save the NM structure at critical points of the analysis: 1) After data entry; 2) after structure initialization; and 3) after preprocessing or analysis. This ensures that no data is lost.

4.14 Clear NeuroMiner structure

This option will clear the NeuroMiner structure **from the workspace**. This will allow the user to select another structure or enter different data.

4.15 Change Working Directory

NeuroMiner will save files in the directory that it is loaded from. This option allows the user to change the working directory.

4.16 Open Manual

This feature is currently not functional. Please open the manual normally.

4.17 Utilities

4.17.0.9 1: Set root paths of neuroimaging tools (SPM/FreeSurfer) If you want to process structural MRI (nii) or surface files from FreeSurfer (mgz) then the paths of these tools need to be established. This option will establish the paths. If the user does not establish the paths, then the neuroimaging functions will not be displayed.

4.17.0.10 2: Create PreprocData path master The option allows the user to update the existing path for only the pre-processed data.

4.17.0.11 3: Create CVdatamat path master This option facilitates the user to update path for the cross-validated data.

4.17.0.12 4: Create CVresults path master This enables to user to configure the cross-validated results in a new directory.

4.17.0.13 5: Create VISdatamat path master This option sets up the user defined-updated path after running the visualisation.

4.17.0.14 6: Create OptPreprocParam path master This is to create a master file that can be used for batch scripts or when repetitive actions are carried out on a set of files the user creates the file list once and does not waste to search specific file sets

4.17.0.15 7: Create OptModel path master this is to create a master file that can be used for batch scripts or when repetitive actions are carried out on a set of files the user creates the file list once and does not waste to search specific file sets.

4.17.0.16 8: Update analyses' paths to new root directory This is an important utility option which allows the users to update the NM structure path, if the user wants to continue the analysis in another directory. By selecting this option, all the paths in the NM structure are updated according to the new/selected location.

5 Example

For a first-time user, NeuroMiner might be best described using an example. This section will run through a simple example using structural neuroimaging data.

5.0.1 Example: Structural NeuroImaging Data

The sample data for this example is open source and can be found in ([link](#)). We processed the data with the Voxel Based Morphometry (VBM) toolbox and we're going to use the `mwrp1` (grey-matter segmented, realigned, warped, modulated) images. Each voxel of these images represents the relative volume of grey matter at any location. The aim of the analysis is to predict who is an individual with a diagnosis of schizophrenia and who doesn't have a diagnosis purely from the grey matter volumes. Based on previous research, the hypothesis is that those with schizophrenia will be able to be distinguished with approximately 70% accuracy.

If the SPM path has been added to NeuroMiner, then the main menu will look like the following:

- 1 : Load data for model discovery and cross-validation
- 2 : Load NeuroMiner structure
- 3 : Change working directory
- 4 : Utilities

The first step is to enter the data using the first option, which will display the following menu:

- 1 : Define machine learning framework [classification]
- 2 : Select data input format [nifti]
- 3 : Define no. of samples and sample identifiers [???]
- 4 : Involve unlabeled cases in the analyses [no]
- 5 : Select space-defining image [?]
- 6 : Describe data [?]

So, we're going to leave the machine learning framework as classification because this is a classification problem, and since we haven't conducted an SPM analysis we will also opt to import NIFTI files directly. I'll then select the third option to define the number of samples and will define the following:

```
Number of groups? : 2
Define name of group number 1 : SCZ
Define name of group number 2 : CTRL
```

I've entered the individuals which a schizophrenia diagnosis first because then the sensitivity accuracy of the analysis will be based on this group rather than the controls. Once this data is entered, then the menu changes to the following:

- 1 : Define machine learning framework [classification]

- 2 : Select data input format [nifti]
- 3 : Modify no. samples [N=2] and sample identifiers [SCZ, CTRL]
- 4 : Involve unlabeled cases in the analyses [no]
- 5 : Select space-defining image [?]
- 6 : Map image files to samples [?]
- 7 : Describe data [?]

I don't want to "Involve unlabeled cases in the analysis", so I'll then select the space-defining image (i.e., the mask; keep in-mind that the images will be resliced to this). I'll either be able to select a region from the WFU PickAtlas or just use my own mask. For this analysis, I don't want to select individual regions, so I'm going to select option 2 and use a pre-prepared ICBM mask in my templates folder. Then because the template is not completely binary and there is a transition from 0 to 1 at the border of the brain, I'm going to select a threshold of 0.5 and then define the regions that are greater than or equal to this threshold (i.e., \geq).

I then select option 6 to "Map image files to samples" and a file selector box will appear that is the same as in SPM. At the top of the box, you will see what the box is for: "Select nifti for sample 1: SCZ". To select the images, I could navigate to the directory on the left and select images, but we recommend to have a file that stores the absolute paths of each of the images in a separate text file. This is because it keeps a record of what has been entered, but mainly because the order of the images is very important for later steps. As an example, my file looks like this:

```
/NeuroDiagStor1/COBRE/Data/0040000/mwrg1mprage.nii
/NeuroDiagStor1/COBRE/Data/0040001/mwrg1mprage.nii
/NeuroDiagStor1/COBRE/Data/0040002/mwrg1mprage.nii
...
```

Once I have this file, I then press the small button on the bottom left called "Ed" and an editor will appear. Then I paste the files into this editor, first for the SCZ group and then for the CTRL group. I then go back to the main menu and enter the global grey matter volumes (or the intracranial volume) with data that I have loaded into the MATLAB workspace. It is critical that these global volume data are in the same order as the subjects that I have just entered into NeuroMiner. If they are in the wrong order then the results of the analysis will be invalid.

I then go back to the main menu and there will now be an option to "Inspect images". I press this, check that the images have been read correctly and then enter a few images to visually inspect. In this case, 1:4. This will display the space-defining mask image as well as the first 4 of my mwrg images. They look good, so I will then choose "Describe data" and enter: COBREmwrg1. After this, I will "IMPORT" the images and this will lead to reslicing and masking with the space-defining image (i.e., my ICBM template).

After the data has been imported, a "MODALITIES" menu will appear with the analysis descriptor (i.e., COBREmwrp1) at the top, then some details about the images. It is important to check that these settings are correct, and the dimensionality of the images. NeuroMiner automatically reports the number of features and the percentage of zero/NaN ("not a number") features. This is important for later analysis (e.g., if there are NaN features then they probably need to be excluded). It is also important to note that for this analysis, no covariates have been found in the NM workspace because they have not been entered yet.

I want to enter covariates of age and sex, so I will choose option "5: Add covariate(s) to NM workspace". These covariates must be already loaded prior to entering into NeuroMiner, otherwise you will have to exit the program, load the covariates, and then enter the program again. After pressing 5, I will then enter the variable that I've previously loaded "cobre_covars" and then name the covariates "one-by-one" because there are only two: age and sex. It's important to note that these variables will just be entered to the NM structure and the **will not be controlled for these variables** until you add this as a preprocessing step.

So now that my variables are in then I'll choose "6: Finish data import for discovery & cross-validation analysis". This means that the data will be saved in the NeuroMiner structure and I don't be able to modify it or add to it again. NeuroMiner locks the data down like this in order to make sure that latter analyses and organisation, relate to the same dataset. So, now that the data are in, I'll be taken to the main NeuroMiner menu where I can set-up my analyses.

- 1 : Inspect data used for model discovery and cross-validation
- 2 : Set up NM parameter workspace
- 3 : Initialize & manage analyses
- 4 : Preprocess features
- 5 : Train supervised classifiers
- 6 : Load NeuroMiner structure
- 7 : Save NeuroMiner structure
- 8 : Clear NeuroMiner structure
- 9 : Change working directory
- 10 : Utilities

As described in section 4, I will need to first set up a parameter template with the settings that I want, initialize the analysis, and then run it. For this reason, I'll go to "2: Set up NM parameter workspace" and it will take me to the following menu (see section 4.2):

- 1 : Cross-validation settings [???]
- 2 : Preprocessing pipeline [...]
- 3 : Classification algorithm [...]
- 4 : Learning algorithm parameters [...]
- 5 : Ensemble generation strategies [...]

- 6 : Visualization options [...]
- 7 : Model saving options [???]
- 8 : Save parameter template
- 9 : Load training template
- 10 : Define verbosity level [Detailed output]

The fields with question marks "???" need to be completed before the analysis can be initialized or run, whereas the other sections are pre-set with default settings. The first thing that I need to do is to set-up cross-validation in step 1. Pressing this enters the cross-validation menu (see section 4.2.2):

- 1 : Select cross-validation framework [(Pooled) cross-validation]
- 2 : Define no. of CV2 permutations [P2 = 10]
- 3 : Define no. of CV2 folds [K2 = 10]
- 4 : Define no. of CV1 permutations [P1 = 10]
- 5 : Define no. of CV1 folds [K1 = 10]
- 6 : Equalize class sizes at the CV1 cycle by undersampling [no]
- 7 : Build CV2/CV1 structure
- 8 : Load CV structure
- 9 : Save CV structure

It's automatically set-up with repeated, nested, cross-validation, which is the gold standard cross-validation set-up currently (see Fig. 5). I could change these settings, but in this case I think they're great for my sample so I'll just choose the option to "7: Build CV2/CV1 structure". This will then build the structure and then I'll go back to the parameter settings menu by simply pressing the enter key because here the default option is to quit the menu.

I'll then enter the pre-processing menu by typing "2" and hitting enter. This will show me a few default settings:

```
=====
CV-PREPROCESSING PIPELINE
=====

>> Step 1: Scale [ from 0 to 1 ], zero-out completely non-finite features
Step 2: Prune non-informative columns from matrix [ Zero Var, Nan, Inf ]
```

And the other settings that I could add are:

- 1 : Enable spatial operations using Spatial OP Wizard
- 2 : Add preprocessing step
- 3 : Remove preprocessing step
- 4 : Insert preprocessing step
- 5 : Replace current preprocessing step
- 6 : Modify current preprocessing step
- 7 : Next step >>
- 8 : Change order of preprocessing steps

For this analysis, I don't want to scale the data first because in previous tests I found that this introduced some noise in the ventricles and I don't need to

prune non-informative columns because I know that there are no NaN or Zero variance features in the data from when I did the import. So I'll remove both of these by pressing "3" and then "3" again.

Now that the CV-PREPROCESSING PIPELINE is empty, I'll add the operations I want using option "2: Add preprocessing step" and will see this menu:

- 1 : Correct data for nuisance covariates using partial correlations
- 2 : Apply dimensionality reduction method to data
- 3 : Standardize data
- 4 : Scale data
- 5 : Normalize to group mean
- 6 : Normalize data to unit vector
- 7 : Apply binning method to data
- 8 : Prune non-informative columns from data matrix
- 9 : Rank / Weight features
- 10 : Correct group offsets from global group mean
- 11 : Extract variance components from data

Then I'll add the option to "Correct data for nuisance covariates" (section 4.2.3.3). I'll then select option 1 and select both covariates by entering "1:2", and I'll keep the rest of the default settings so that age and sex will be regressed from the entire sample. After I've set these options, I'll then quit to go back to the preprocessing menu.

After covarying for nuisance covariates, I now want to perform a dimensionality reduction so I will select "2: add preprocessing step" and then "Apply dimensionality reduction method to data" (section 4.2.3.4). There are a lot of options, but basically it seems that PCA works well most of the time so I'll select "1: Principal components analysis". Then I'll keep all the defaults because these work well too, and return to the main menu. In this step, if I had have wanted to optimise over a number of different component reductions, then I could have defined the energy as 1, returned to the preprocessing menu, added another step, and then added extracted component subspaces and defined something like [0.2 0.4 0.6 0.8] to see how the algorithm performs when I retain from 20 to 80% of the energy.

Then I want to add some scaling because I'm going to use an SVM and the data must be scaled for this and also because it makes the results more interpretable. So I'll select "4: Scale data" and again just accept the defaults. Now the CV PREPROCESSING PIPELINE should include:

```
=====
CV-PREPROCESSING PIPELINE
=====

>> Step 1: Correct for nuisance covariates [age,sex]
Step 2: Apply dimensionality reduction [ PCA ]
Step 3: Scale [ from 0 to 1 ], zero-out completely non-finite features
```

And this will be applied within each of the CV1 folds.

I'll then go back to the main parameter template menu by pressing enter and will see this:

- 1 : Cross-validation settings [...]
- 2 : Preprocessing pipeline [...]
- 3 : Classification algorithm [...]
- 4 : Learning algorithm parameters [...]
- 5 : Ensemble generation strategies [...]
- 6 : Visualization options [...]
- 7 : Model saving options [???]
- 8 : Save parameter template
- 9 : Load training template
- 10 : Define verbosity level [Detailed output]

Then I want to define the classification algorithm by pressing "3: Classification algorithm" (see section [4.2.4](#)). The default in NeuroMiner is to implement a SVM, and the only thing that I want to change is the performance measure to Balanced Accuracy instead of accuracy. I can do this by selecting the "performance criterion" and then selecting "Balanced accuracy: BAC...". Then I'll return to the parameter template menu by pressing enter again.

Here I could change the learning algorithm parameters (see section [4.2.5](#)), but I want to keep the defaults again. It's important to note here that I'm keeping a range of C parameters that will be optimised during the nested cross-validation. I could also enable some filters and wrappers (see section [4.2.6](#)), but these are more advanced features that can be addressed in a further tutorial. I could also enable visualization options (such as the derivation of Z or p-values; see section [4.2.7](#)), but because I have enabled PCA these will be overly conservative and so I won't use them.

The final thing to do is to specify "7: Model saving options". For this, I'll choose that I don't want to save the models because it will take up a lot of space (i.e., more than 10 000 models). Then I'll define a name anyway because this is required in NeuroMiner. Then I'll return to the MAIN INTERFACE menu of NeuroMiner by pressing the enter key until I get back here:

- 1 : Inspect data used for model discovery and cross-validation
- 2 : Set up NM parameter workspace
- 3 : Initialize & manage analyses
- 4 : Preprocess features
- 5 : Train supervised classifiers
- 6 : Load NeuroMiner structure
- 7 : Save NeuroMiner structure
- 8 : Clear NeuroMiner structure
- 9 : Change working directory
- 10 : Utilities

From this menu, I'll then need to "3: initialize & manage analyses" (see section [4.5](#)) where I'll see this:

```
Define analysis identifier [ ? ]
Provide analysis description [ ? ]
Specify parent directory of the analysis [ ? ]
PROCEED >>>
```

I'll then define the identifier as "COBRE_SVM_analysis1", provide a description of the analysis "This is a test analysis using the COBRE dataset", and specify the parent directory of the analysis using the file selector box that pops up. Then I'll "PROCEED" and a directory will be created in the directory. Within this directory you will find a log file that outlines the operating system and the settings that I have just defined. After assuring myself that these are ok, then I'll go back to the main menu.

Now I have the option to "4: preprocess features" (see section [4.6](#)) first, which will apply all the preprocessing steps I defined to the CV1 folds. Or alternatively, I could simply "5: Train supervised classifiers" (see section [4.7](#) from scratch and automatically preprocess the features as part of this. I want to preprocess first here because then I can run other parameters later with this preprocessed data, so I'll select "4: Preprocess features". Which will activate:

```
1 : Overwrite existing preprocessed data files [ no ]
2 : Use existing preprocessing parameter files, if available [ no ]
3 : Write out computed parameters to disk (may require A LOT of disk space)
[ no ]
4 : Select CV2 partitions to operate on [ 100 selected ]
5 : PROCEED >>>
```

I'll keep all the defaults here, but I'll select "4: Select CV2 partitions to operate on", then hit "4: Deselect all CV2 positions", and then I'll choose "3: Select single perm / fold range" and enter [1,1,1,10] to select the first permutation and folds 1-10. This is because selecting all permutations will take a lot of time and is not required for this example. Then I'll hit enter to go back to the main menu and then "5: PROCEED". The analysis should run now, and I'll go and get some lunch.

When I get back, the analysis should have finished and NeuroMiner should have returned to the main menu. All the preprocessed data files should be stored in the subdirectory for this analysis that I activated previously. Now all that's left to do is to actually train the classifiers. This can be done by selecting "5: Train supervised classifiers" then "2: Operation mode of ML training module" and select "2: Compute using existing preprocdata-MATs" and when the file selector box appears, I'll select the preprocessed datamats from the previous analysis. Then I'll select "PROCEED", check to get an initial impression of the accuracy , and then I'll go and write that fellowship application.

When I come back again, the analysis is finished and the training datamats

are stored in the directory as well. I can now see the results by going to "7: Display the training results" (see [4.9](#)), including the accuracies (see [7](#)). This is interesting and so now I want to visualise the weights on the brain. For this, I'll go back to the main menu and select "Visualise Classifiers" (see section [4.8](#)) and select the data that I've processed and hit proceed. This will then create some images that I can view with a viewer such as MRICRON, FSL, SPM, MANGO, or others.

If this is your own data, hopefully then you have found out about some problem and you can now go and publish.

6 Known Issues

NeuroMiner is currently undergoing testing and there are a number of known issues that require further work. These are the following.

- The results of external validation analyses are not displayed and need to be accessed in the NM structure;
- External validation analysis currently uses a lot of RAM for high dimensional matrices;
- External validation does not work when using stacking;
- If training is run without preprocessing then it can use a lot of RAM for high dimensional matrices;
- Random forest algorithm is not included in precompiled Linux versions and needs to be compiled;
- There is currently no facility to visualise surface-based (e.g., mgz) results and it is recommended that this is achieved manually;
- GLMNET is unstable under Windows 7
- The database import does not work with Linux odt files. Please convert to xls.
- The extraction of scores (e.g., decision scores) in the results viewer does not function using Linux